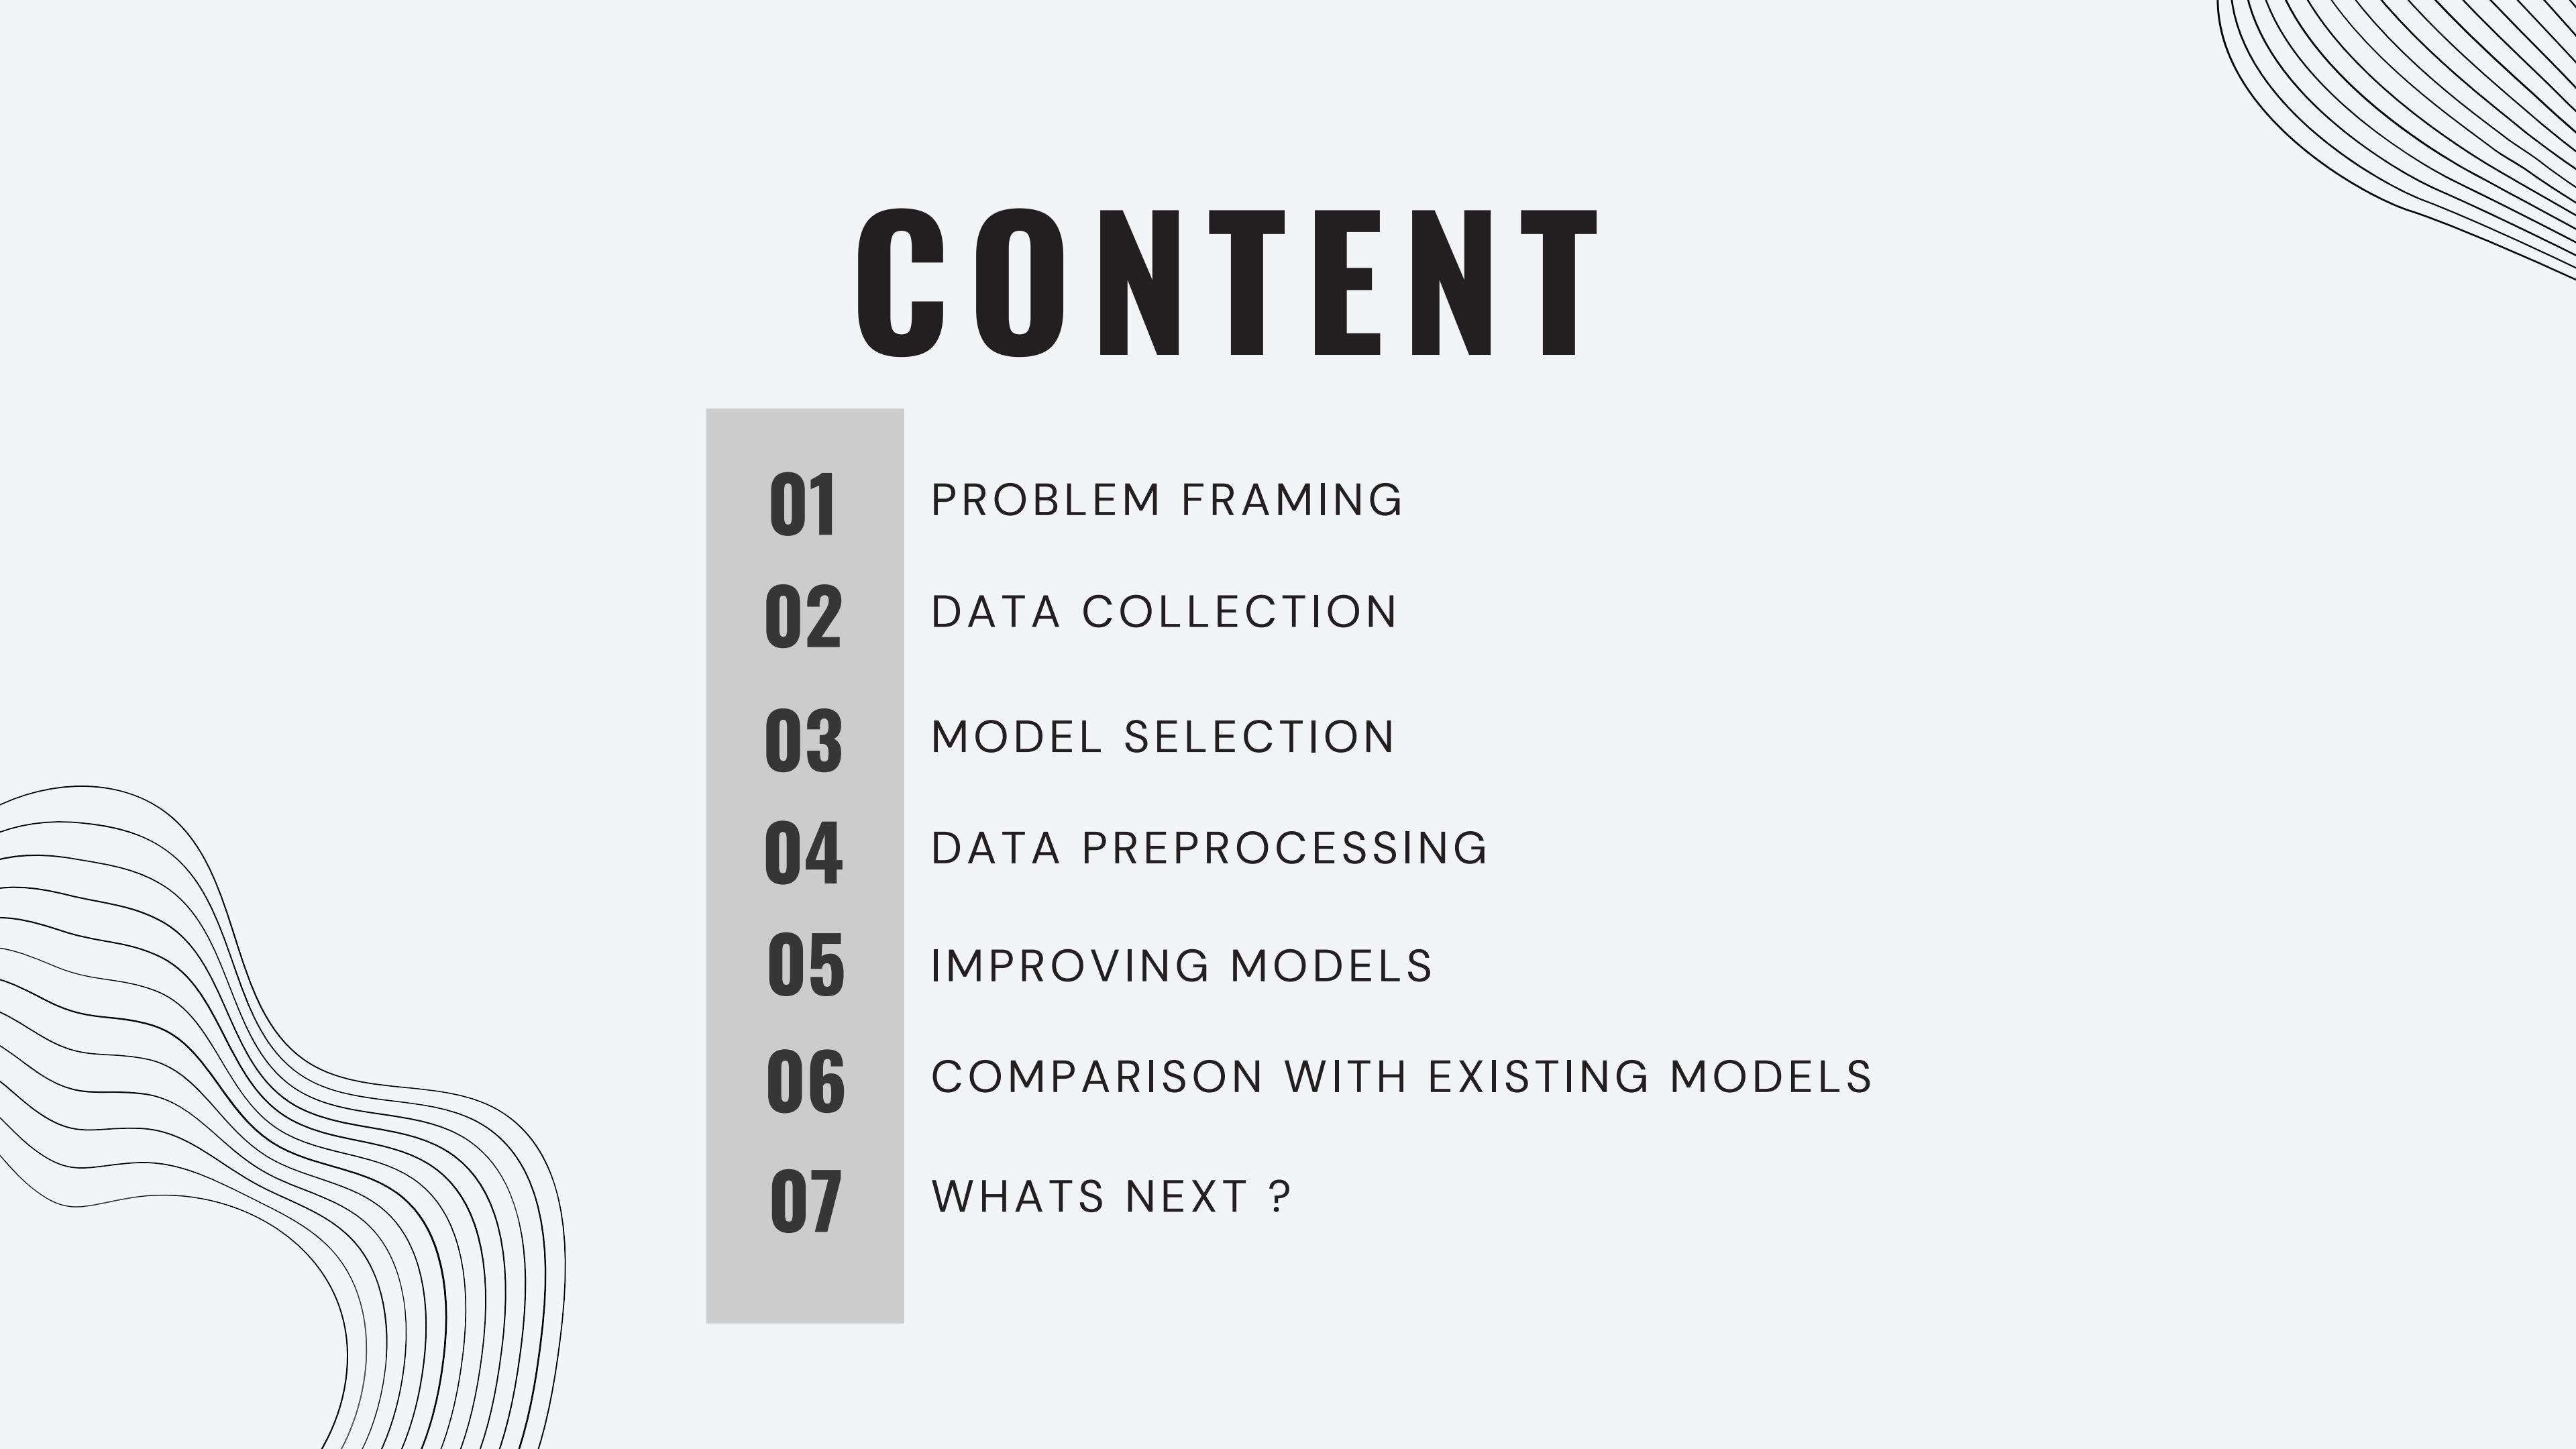


AI PROJECT

EXPLORE TRANSFER LEARNING THROUGH SMALL AND LARGE MODELS

BY: IKRAM ACHALI / LOKESHWARN VENGADABADY

CONTENT

- 
- 01** PROBLEM FRAMING
 - 02** DATA COLLECTION
 - 03** MODEL SELECTION
 - 04** DATA PREPROCESSING
 - 05** IMPROVING MODELS
 - 06** COMPARISON WITH EXISTING MODELS
 - 07** WHATS NEXT ?

TRANSFER LEARNING

- Pre-Trained Knowledge Reuse
- Fine-Tuning
- Efficiency and Performance



DATA COLLECTION

7,591 images (of age 1-100) with associated real and apparent age labels

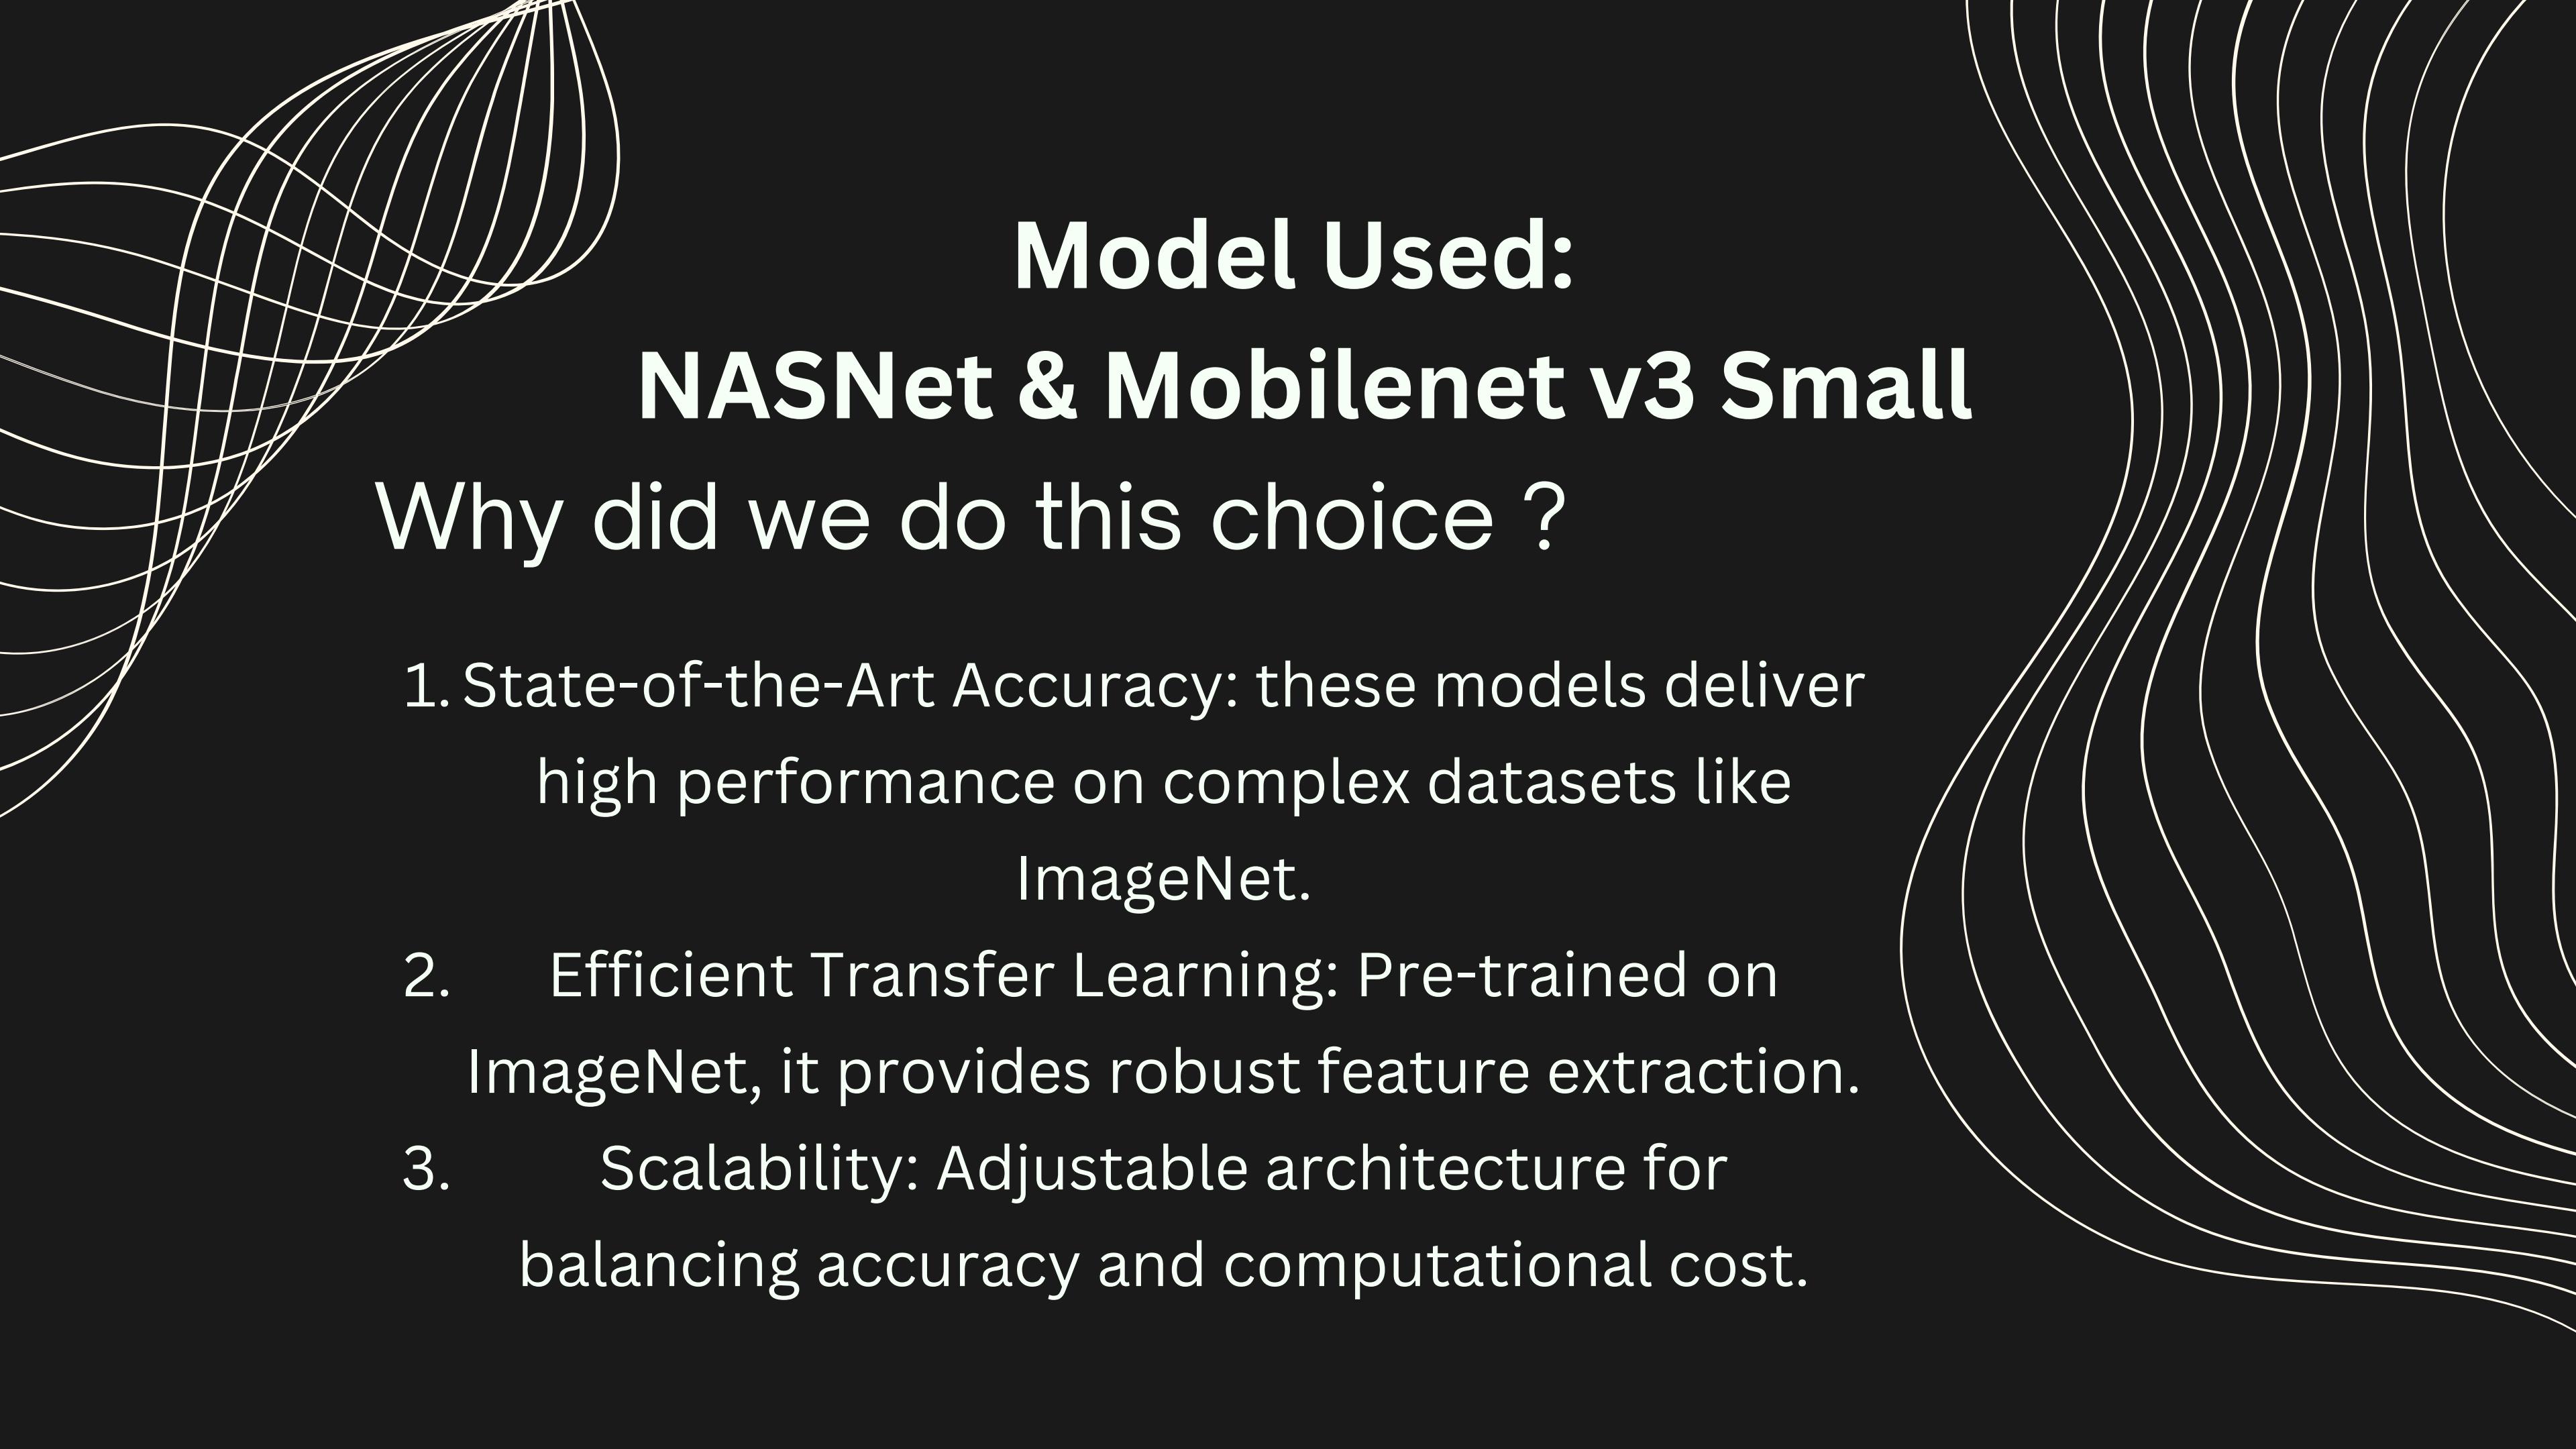
**APPA-REAL
DATABASE**



Ex 1:
age : 1



Ex 2 :
age : 60



Model Used:

NASNet & Mobilenet v3 Small

Why did we do this choice ?

1. State-of-the-Art Accuracy: these models deliver high performance on complex datasets like ImageNet.
2. Efficient Transfer Learning: Pre-trained on ImageNet, it provides robust feature extraction.
3. Scalability: Adjustable architecture for balancing accuracy and computational cost.



NASNET
TRANSFER

**LET'S DISCOVER THE
PROCESS**

Preprocessing IMAGES

PREPROCESSING

7,591 images
with associated
real and
apparent age
labels

**APPA-REAL
DATABASE**

Read images in
the dataset and
save them in
array format in
the npy files

**IMAGES TO
NPY FILES**

Image arrays to
be between 0
and 255

**READY TO BE
USED**

```
### Train CSV ###
```

	file_name	num_ratings	apparent_age_avg	apparent_age_std	real_age
0	000000.jpg	36	5.000000	1.146423	4
1	000001.jpg	63	20.079365	4.096819	18
2	000002.jpg	38	76.815789	6.133009	80
3	000003.jpg	38	55.657895	7.864653	50
4	000004.jpg	15	17.666667	3.457222	17

```
### Test CSV ###
```

	file_name	num_ratings	apparent_age_avg	apparent_age_std	real_age
0	005613.jpg	39	23.205128	5.530678	19
1	005614.jpg	38	70.736842	6.570549	76
2	005615.jpg	38	55.368421	5.874457	40
3	005616.jpg	36	24.277778	2.224681	21
4	005617.jpg	39	25.230769	5.691460	34

```
### Valid CSV ###
```

	file_name	num_ratings	apparent_age_avg	apparent_age_std	real_age
0	004113.jpg	26	26.230769	4.003076	29
1	004114.jpg	39	27.256410	4.552142	25
2	004115.jpg	14	23.142857	7.262730	37
3	004116.jpg	38	73.289474	6.472010	80
4	004117.jpg	21	20.142857	2.797958	25

```
if (generate_X == True):
    for i in train_img_names[:Nb_train]:
        img = Image.open(f'{base_dir}/train/{i)').convert("RGB")
        img = img.resize((224,224))
        img_array = np.array(img)
        image_list.append(img_array)
        if (float(i.strip('.jpg')) % 1000==0):
            print(f"completed {i.strip('.jpg')} images")
    np.save(f'{base_dir}/data_train.npy', image_list)
```

```
base_model = tf.keras.applications.NASNetMobile(input_shape=(224, 224, 3), weights='imagenet', include_top=False)
preprocess_input_Nasnet = tf.keras.applications.nasnet.preprocess_input
Nasnet_model = tf.keras.models.Model(inputs=base_model.input, outputs=output)
X_train_processed_Nasnet = preprocess_input_Nasnet(np.copy(X_train))
X_valid_processed_Nasnet = preprocess_input_Nasnet(np.copy(X_valid))

base_model1 = tf.keras.applications.MobileNetV3Small(input_shape=(224, 224, 3), weights='imagenet', include_top=False)
preprocess_input_Mobilenet = tf.keras.applications.mobilenet_v3.preprocess_input
X_train_processed_Mobilenet = preprocess_input_Mobilenet(np.copy(X_train))
X_valid_processed_Mobilenet = preprocess_input_Mobilenet(np.copy(X_valid))
```

Train Model

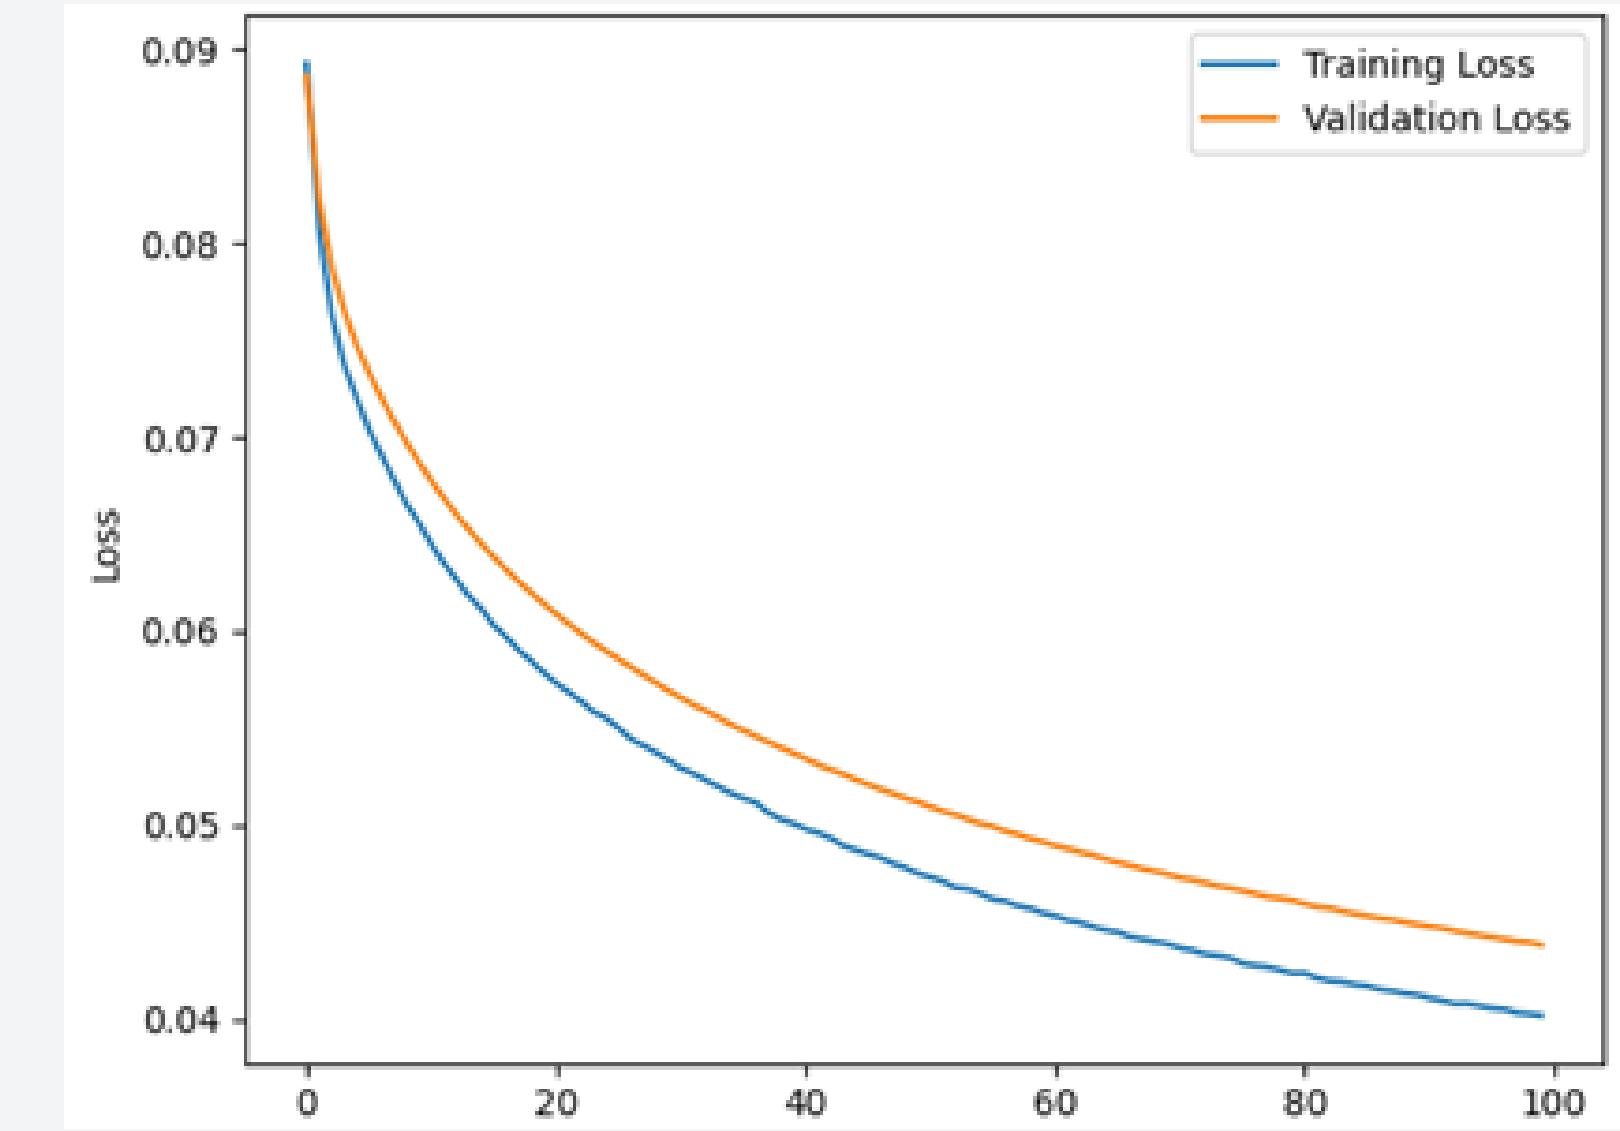
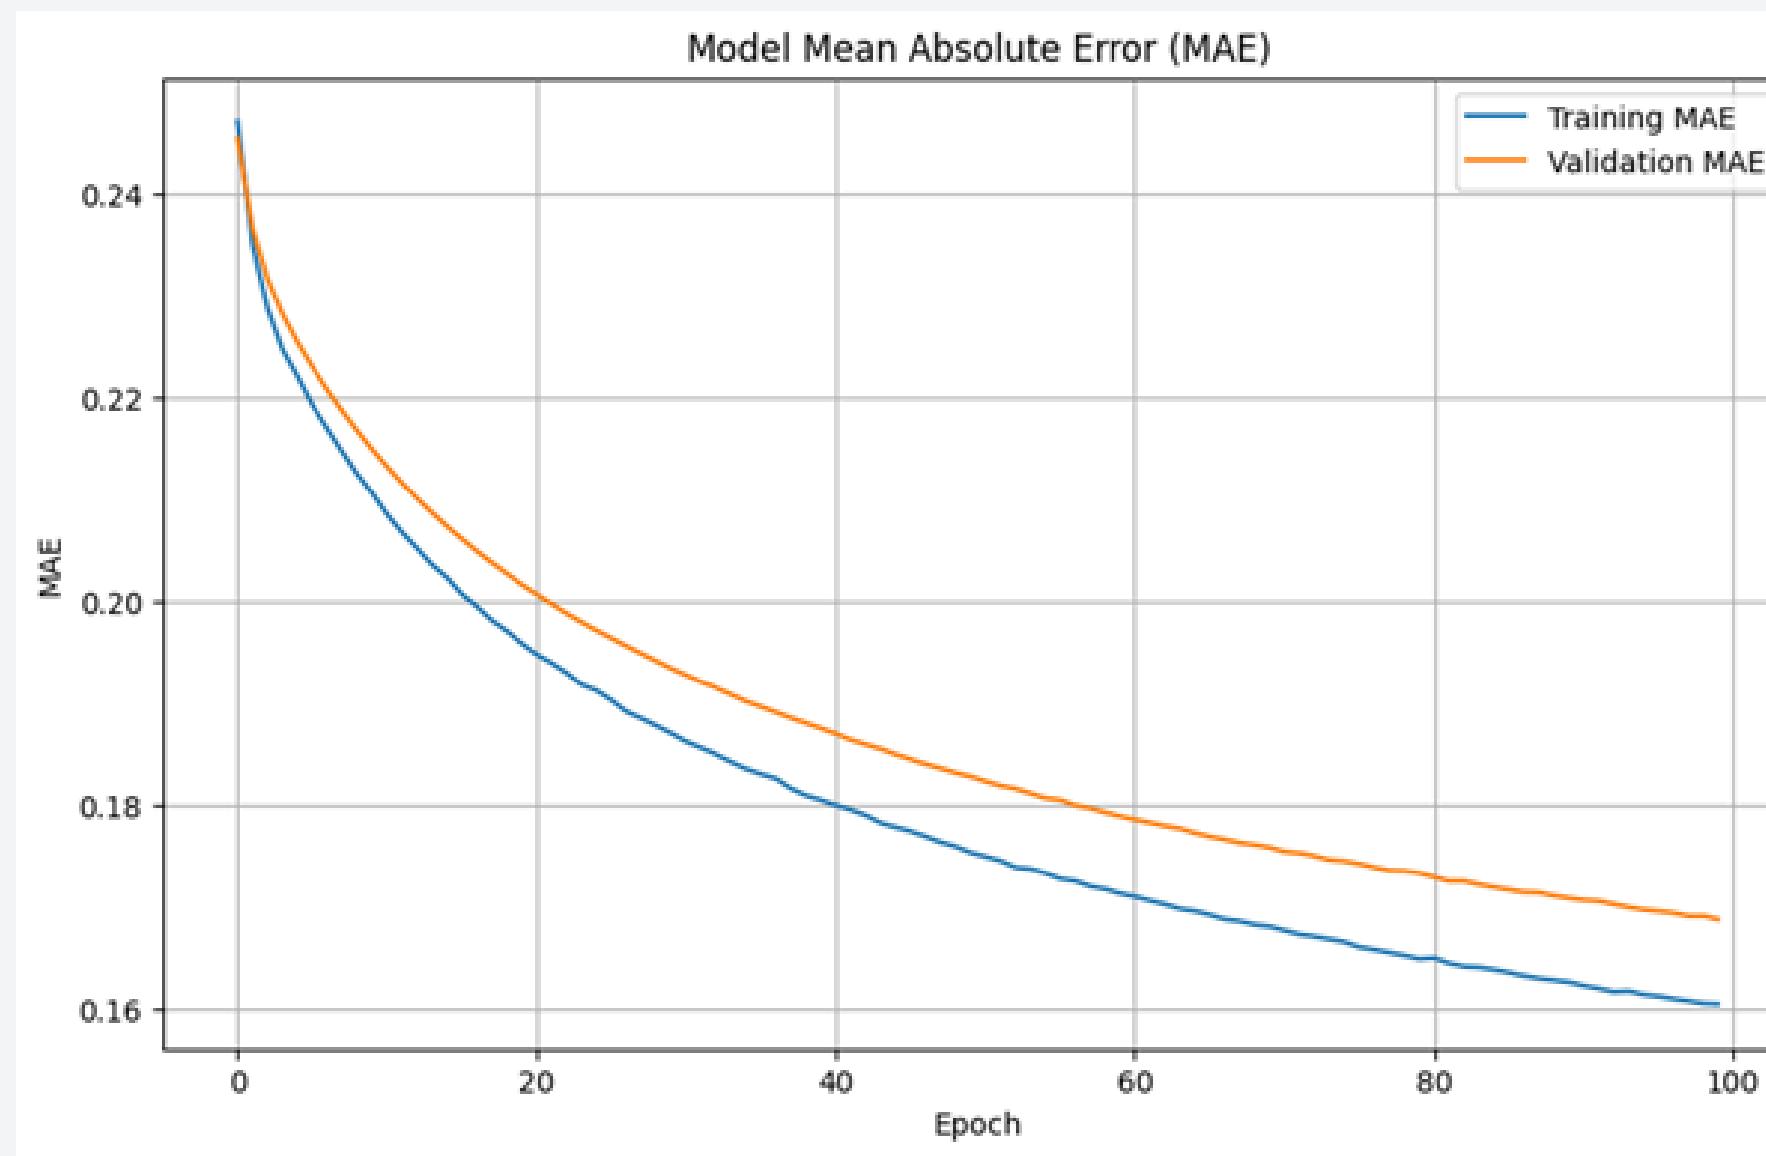
ORIGINAL MODEL + ONE
OUTPUT LAYER

Nasnet

```
[12] ### Nasnet basic
early_stopping_Nasnet = tf.keras.callbacks.EarlyStopping(monitor='val_loss', mode='min', verbose=1, patience=6)
# Load NASNetMobile with pre-trained ImageNet weights, excluding the top layer
base_model = tf.keras.applications.NASNetMobile(input_shape=(224, 224, 3), weights='imagenet', include_top=False)
#preprocess_input = tf.keras.applications.mobilenet.preprocess_input
preprocess_input_Nasnet = tf.keras.applications.nasnet.preprocess_input
base_model.trainable = False
# Adding custom layers
x = base_model.output
output = tf.keras.layers.Dense(1, activation='sigmoid')(x)
# Create the new model with NASNetMobile as the base
Nasnet_model = tf.keras.models.Model(inputs=base_model.input, outputs=output)
# Preprocess the data
X_train_processed_Nasnet = preprocess_input_Nasnet(np.copy(X_train))
X_valid_processed_Nasnet = preprocess_input_Nasnet(np.copy(X_valid))
print("Num GPUs Available: ", len(tf.config.list_physical_devices('GPU')))
```

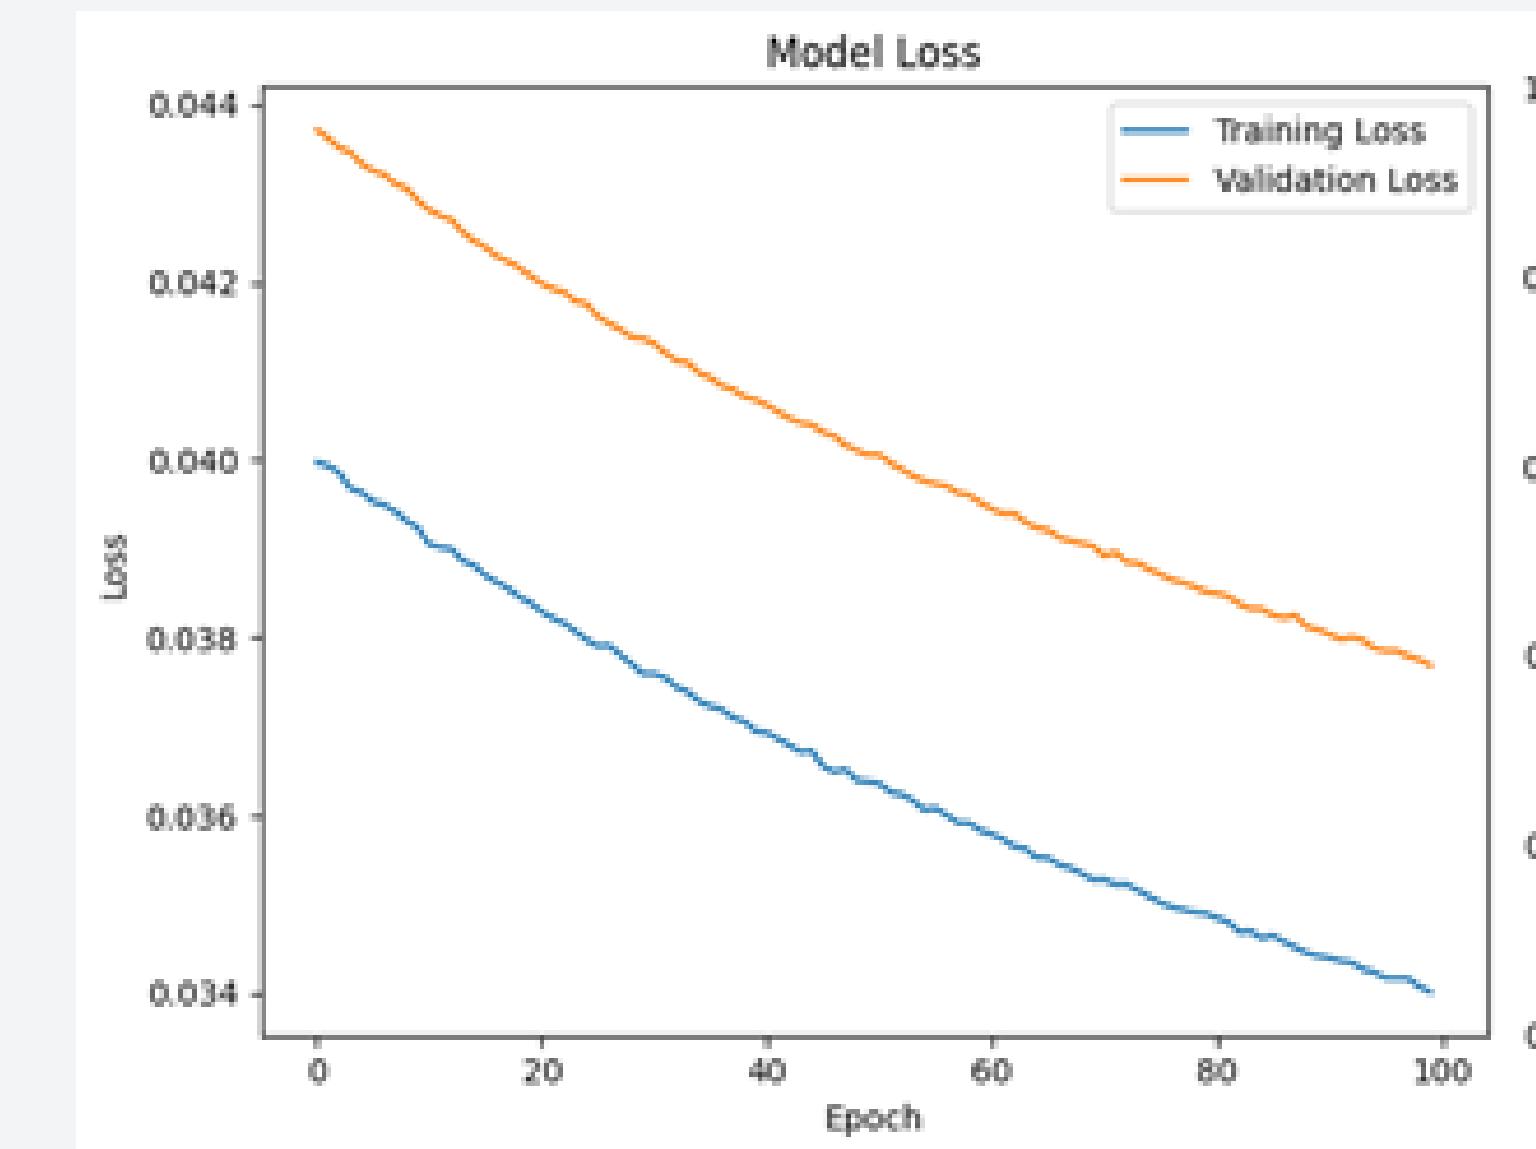
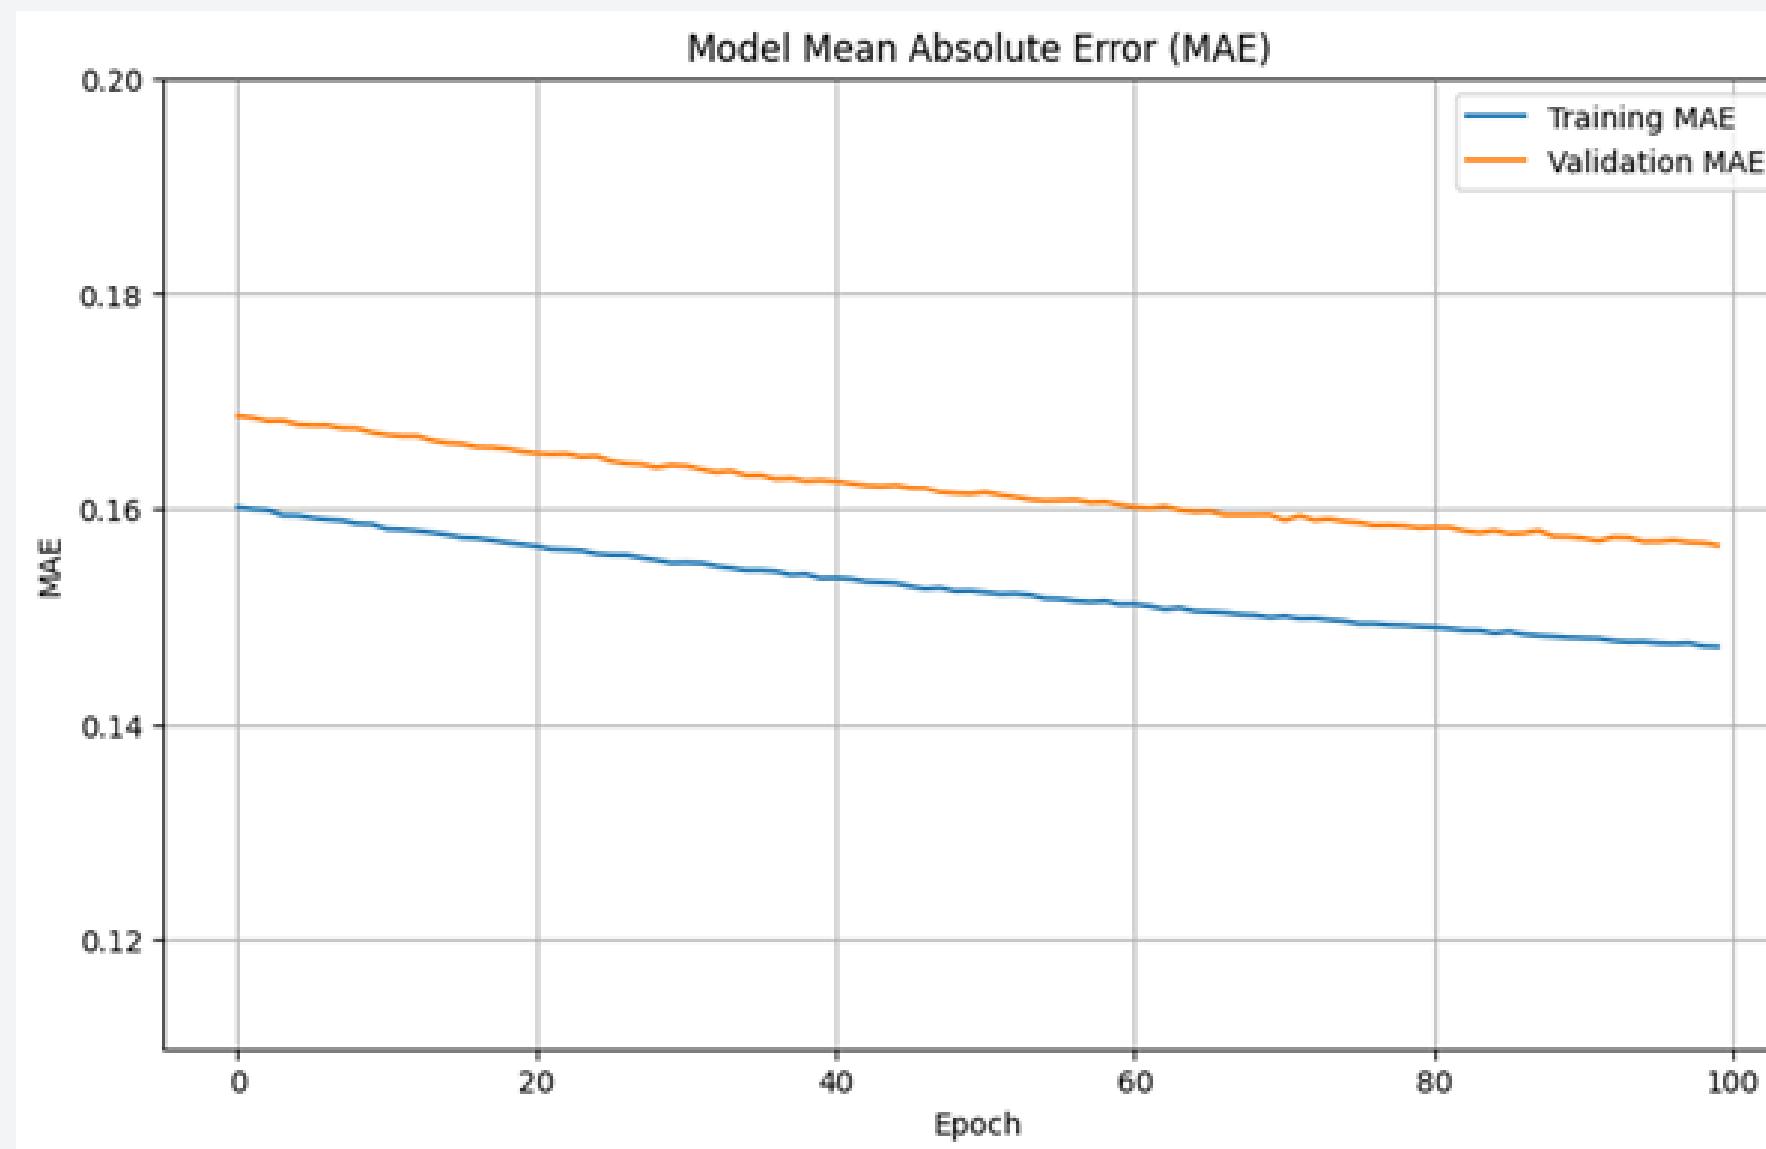
```
→ Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/nasnet/NASNet-mobile-no-top.h5
19993432/19993432 0s 0us/step
Num GPUs Available: 1
```

Nasnet



Nasnet

WITH MORE EPOCH



Mobilenet v3 Small

```
[19] ### Mobilenet basic
    early_stopping_MobileNet = tf.keras.callbacks.EarlyStopping(monitor='val_loss', mode='min', verbose=1, patience=7)

    #base_model = tf.keras.applications.ResNet50(input_shape=(224, 224, 3), weights='imagenet', include_top=False)
    base_model = tf.keras.applications.MobileNetV3Small(input_shape=(224, 224, 3), weights='imagenet', include_top=False)

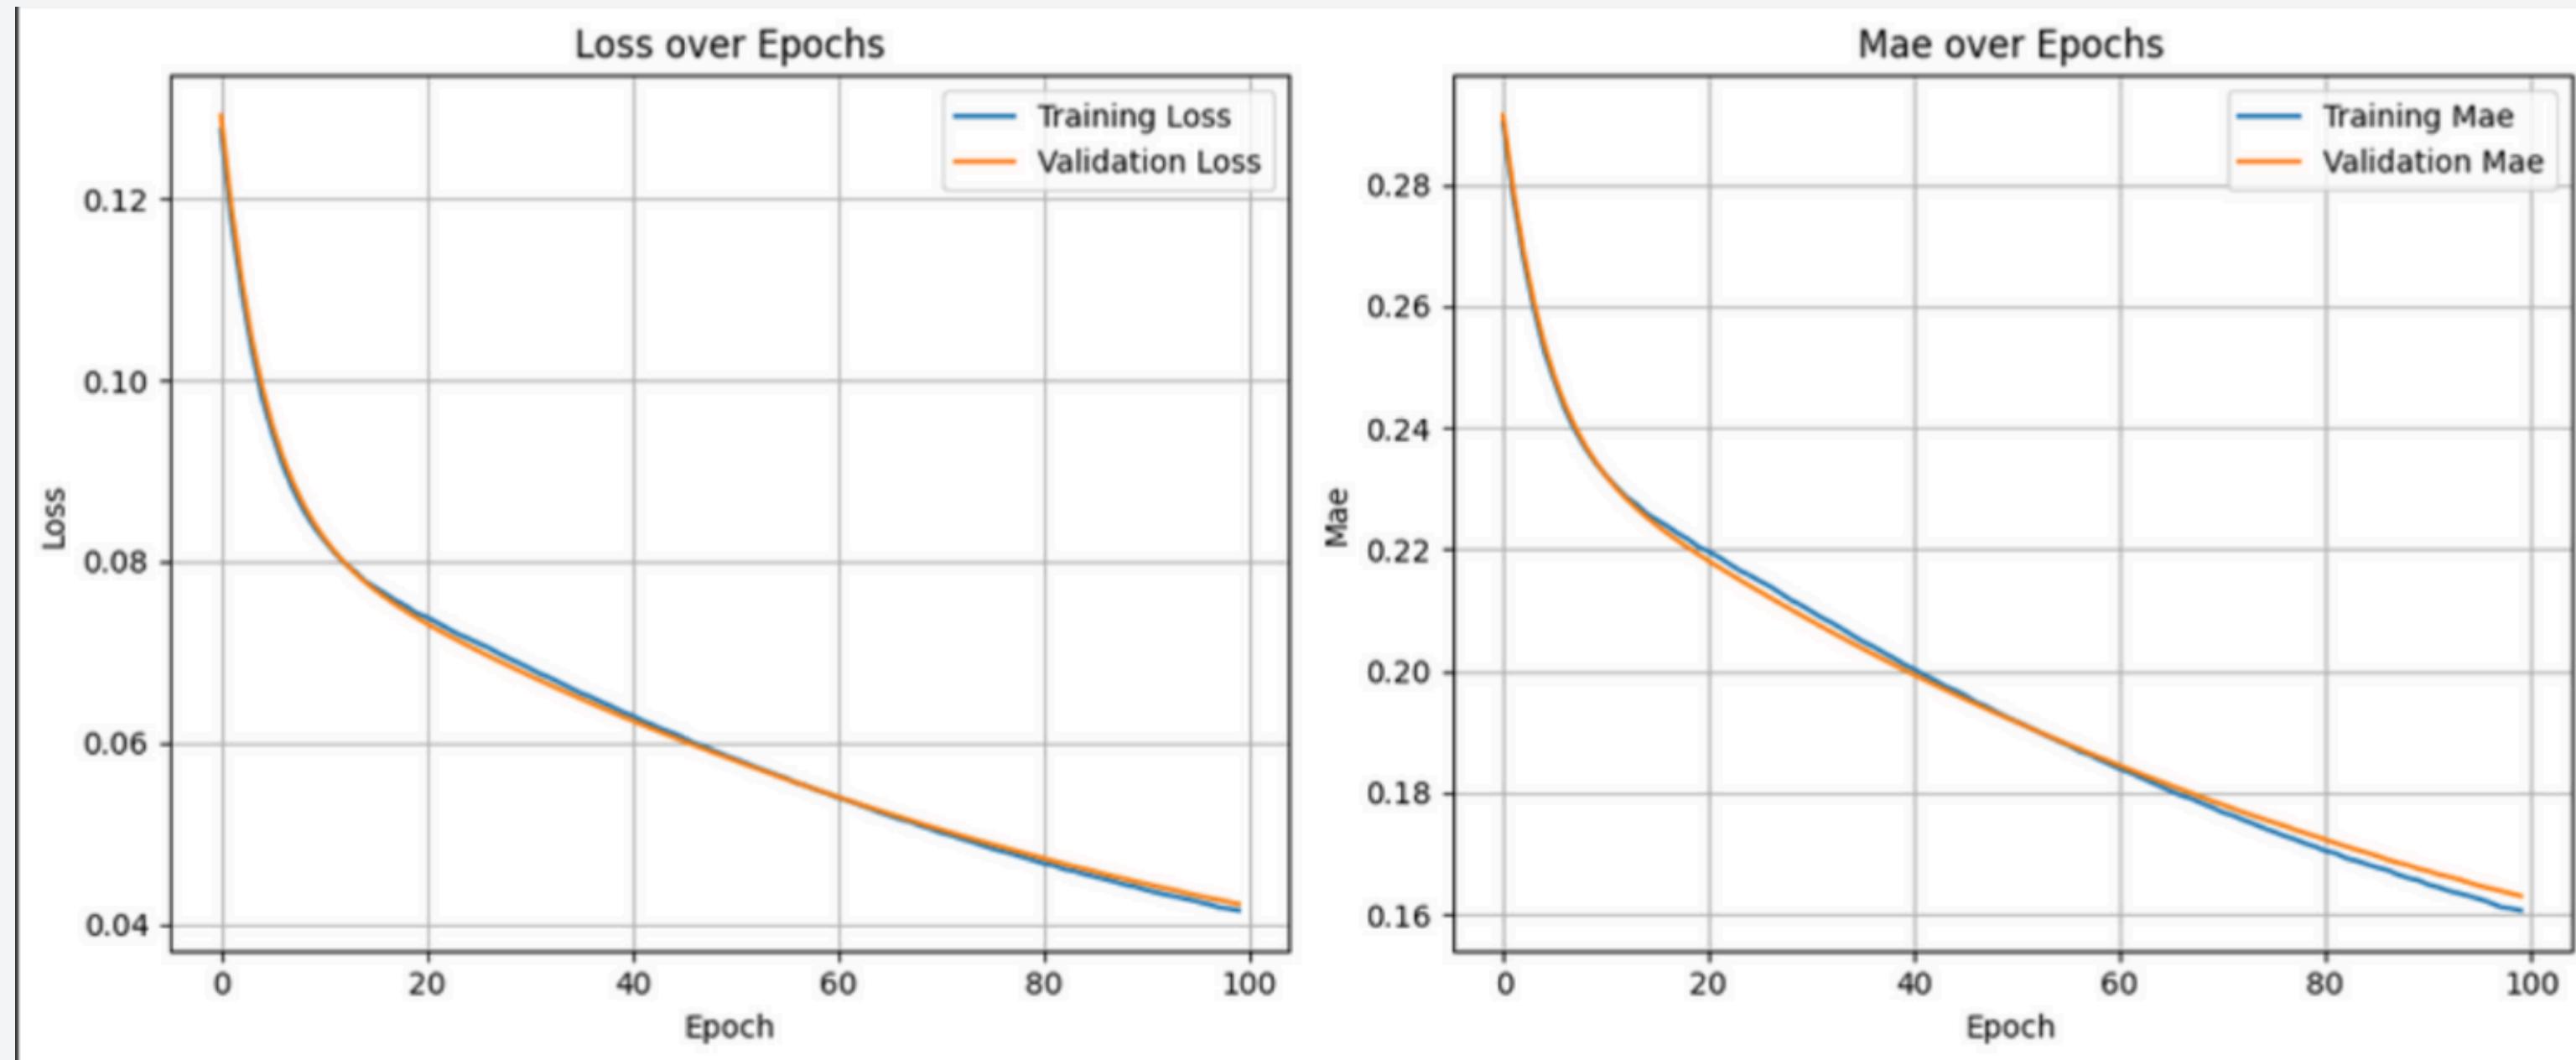
    #preprocess_input = tf.keras.applications.mobilenet.preprocess_input
    preprocess_input_Mobilenet = tf.keras.applications.mobilenet_v3.preprocess_input

    base_model.trainable = False

    # Add custom layers
    #x = tf.keras.layers.GlobalAveragePooling2D()(base_model.output)
    #x = tf.keras.layers.Dense(1024, activation='relu')(x)
    #x = tf.keras.layers.Dropout(0.5)(x)
    output = tf.keras.layers.Dense(1, activation='sigmoid')(base_model.output)

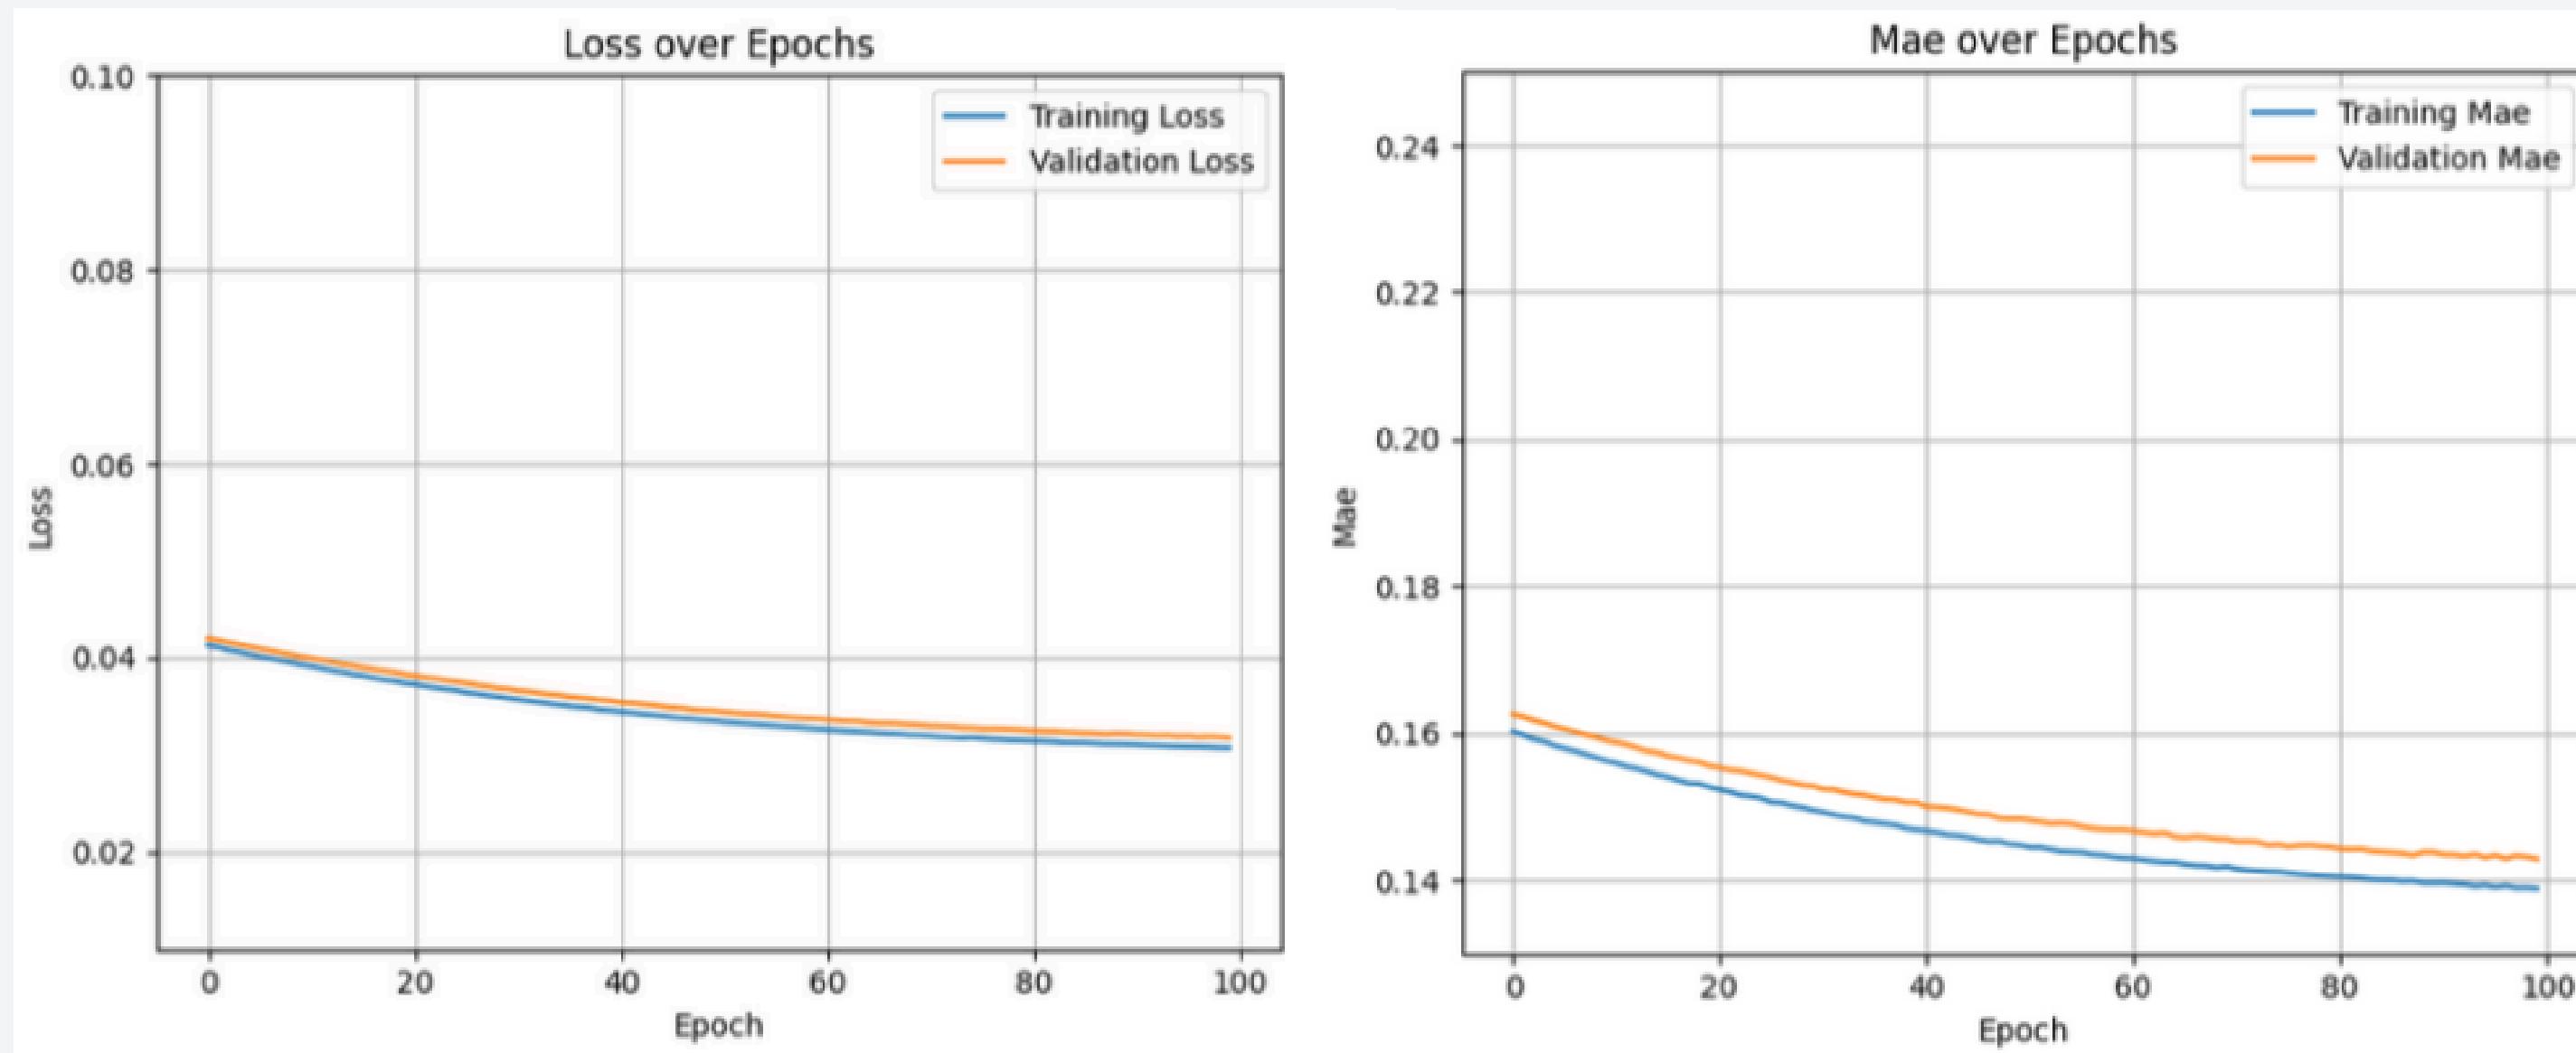
    # Create the new model with MobileNet as the base
    Mobilenet_model = tf.keras.models.Model(inputs=base_model.input, outputs=output)
    # Preprocess the data
    X_train_processed_Mobilenet = preprocess_input_Mobilenet(np.copy(X_train))
    X_valid_processed_Mobilenet = preprocess_input_Mobilenet(np.copy(X_valid))
```

Mobilenet v3 Small



Mobilenet v3 Small

WITH MORE EPOCH



Train Model

UNFREEZED LAYERS + GLOBAL
AVERAGE POOLING & RELU
LAYERS

```

### Nasnet implementation avec 0.7 de Dropout and with 1 trainable layer
early_stopping_Nasnet = tf.keras.callbacks.EarlyStopping(monitor='val_loss', mode='min', verbose=1, patience=6)
# Load NASNetMobile with pre-trained ImageNet weights, excluding the top layer
base_model = tf.keras.applications.NASNetMobile(input_shape=(224, 224, 3), weights='imagenet', include_top=False)
#preprocess_input = tf.keras.applications.mobilenet.preprocess_input
preprocess_input_Nasnet = tf.keras.applications.nasnet.preprocess_input
base_model.trainable = False
for layer in base_model.layers[-1:]:
    layer.trainable = True
# Adding custom layers
x = tf.keras.layers.GlobalAveragePooling2D()(base_model.output)
x = tf.keras.layers.Dense(1024, activation='relu')(x)
x = tf.keras.layers.Dropout(0.7)(x) # Dropout with 50% probability
output = tf.keras.layers.Dense(1, activation='sigmoid')(x)
# Create the new model with NASNetMobile as the base
Nasnet_model = tf.keras.models.Model(inputs=base_model.input, outputs=output)
# Preprocess the data
X_train_processed_Nasnet = preprocess_input_Nasnet(np.copy(X_train))
X_valid_processed_Nasnet = preprocess_input_Nasnet(np.copy(X_valid))
print("Num GPUs Available: ", len(tf.config.list_physical_devices('GPU')))


```

```

Epoch 47/100
125/125 ━━━━━━━━━━━━ 4s 34ms/step - loss: 0.0156 - mae: 0.0982 - val_loss: 0.0217 - val_mae: 0.1155
Epoch 48/100
125/125 ━━━━━━━━━━━━ 4s 34ms/step - loss: 0.0164 - mae: 0.0999 - val_loss: 0.0216 - val_mae: 0.1160
Epoch 49/100
125/125 ━━━━━━━━━━━━ 4s 34ms/step - loss: 0.0150 - mae: 0.0960 - val_loss: 0.0216 - val_mae: 0.1154
Epoch 50/100
125/125 ━━━━━━━━━━━━ 4s 34ms/step - loss: 0.0143 - mae: 0.0942 - val_loss: 0.0216 - val_mae: 0.1152
Epoch 51/100
125/125 ━━━━━━━━━━━━ 4s 35ms/step - loss: 0.0138 - mae: 0.0926 - val_loss: 0.0216 - val_mae: 0.1151
Epoch 52/100
125/125 ━━━━━━━━━━━━ 4s 35ms/step - loss: 0.0147 - mae: 0.0944 - val_loss: 0.0216 - val_mae: 0.1153
Epoch 53/100
125/125 ━━━━━━━━━━━━ 4s 35ms/step - loss: 0.0139 - mae: 0.0919 - val_loss: 0.0216 - val_mae: 0.1155
Epoch 54/100
125/125 ━━━━━━━━━━━━ 4s 35ms/step - loss: 0.0140 - mae: 0.0933 - val_loss: 0.0216 - val_mae: 0.1157
Epoch 55/100

```

```
[12] ### Mobilenet basic
2s
    early_stopping_MobileNet = tf.keras.callbacks.EarlyStopping(monitor='val_loss', mode='min', verbose=1, patience=7)
#base_model = tf.keras.applications.ResNet50(input_shape=(224, 224, 3), weights='imagenet', include_top=False)
base_model = tf.keras.applications.MobileNetV3Small(input_shape=(224, 224, 3), weights='imagenet', include_top=False)
#preprocess_input = tf.keras.applications.mobilenet.preprocess_input
preprocess_input_Mobilenet = tf.keras.applications.mobilenet_v3.preprocess_input
base_model.trainable = False
base_model.trainable = False
for layer in base_model.layers[-1:]:
    layer.trainable = True

# Add custom layers
x = tf.keras.layers.GlobalAveragePooling2D()(base_model.output)
x = tf.keras.layers.Dense(1024, activation='relu')(x)
x = tf.keras.layers.Dropout(0.7)(x)
output = tf.keras.layers.Dense(1, activation='sigmoid')(x)

# Create the new model with MobileNet as the base
Mobilenet_model = tf.keras.models.Model(inputs=base_model.input, outputs=output)
# Preprocess the data
X_train_processed_Mobilenet = preprocess_input_Mobilenet(np.copy(X_train))
X_valid_processed_Mobilenet = preprocess_input_Mobilenet(np.copy(X_valid))
```

```
125/125 ━━━━━━━━━━ 1s 7ms/step - loss: 0.0124 - mae: 0.0876 - val_loss: 0.0176 - val_mae: 0.1028
Epoch 95/100
125/125 ━━━━━━━━━━ 1s 7ms/step - loss: 0.0129 - mae: 0.0894 - val_loss: 0.0176 - val_mae: 0.1030
Epoch 96/100
125/125 ━━━━━━━━━━ 1s 7ms/step - loss: 0.0127 - mae: 0.0898 - val_loss: 0.0176 - val_mae: 0.1028
Epoch 97/100
125/125 ━━━━━━━━━━ 1s 7ms/step - loss: 0.0126 - mae: 0.0883 - val_loss: 0.0176 - val_mae: 0.1027
Epoch 98/100
125/125 ━━━━━━━━━━ 1s 7ms/step - loss: 0.0126 - mae: 0.0888 - val_loss: 0.0176 - val_mae: 0.1032
Epoch 99/100
125/125 ━━━━━━━━━━ 1s 7ms/step - loss: 0.0123 - mae: 0.0880 - val_loss: 0.0176 - val_mae: 0.1031
Epoch 100/100
125/125 ━━━━━━━━━━ 1s 7ms/step - loss: 0.0126 - mae: 0.0878 - val_loss: 0.0176 - val_mae: 0.1029
```

Train Model
UNFREEZE MORE

```
early_stopping_monitor = tf.keras.callbacks.EarlyStopping(monitor='val_loss', mode='min', verbose=1, patience=5)

#base_model = tf.keras.applications.ResNet50(input_shape=(224, 224, 3), weights='imagenet', include_top=False)
base_model = tf.keras.applications.MobileNetV3Small(input_shape=(224, 224, 3), weights='imagenet', include_top=False)

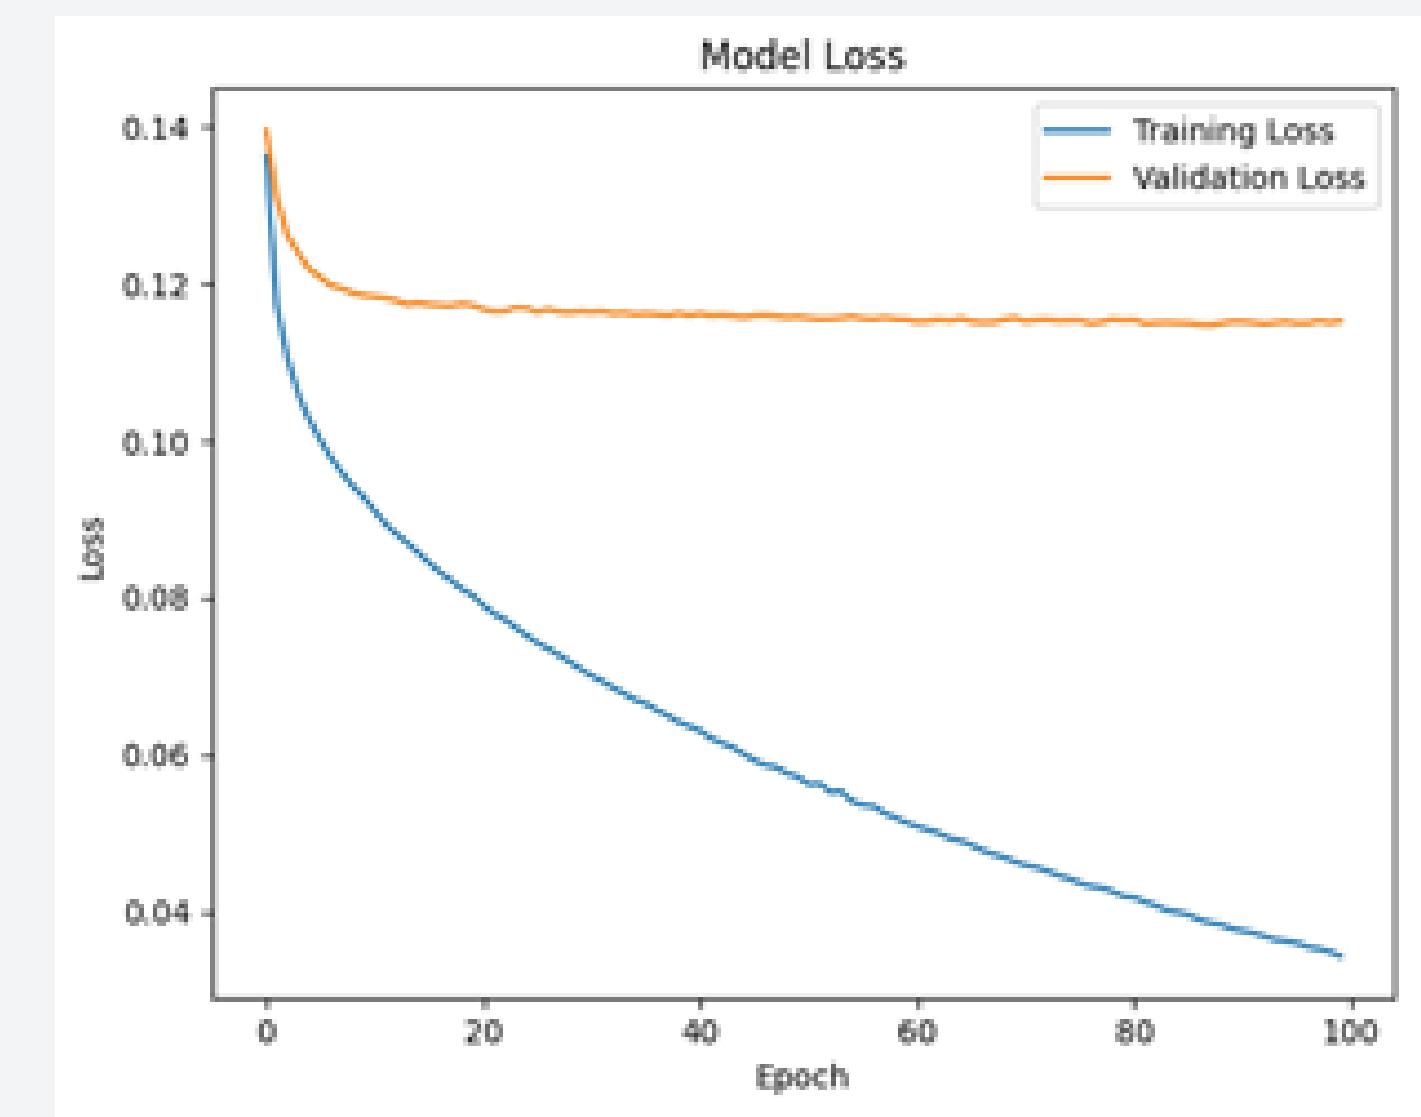
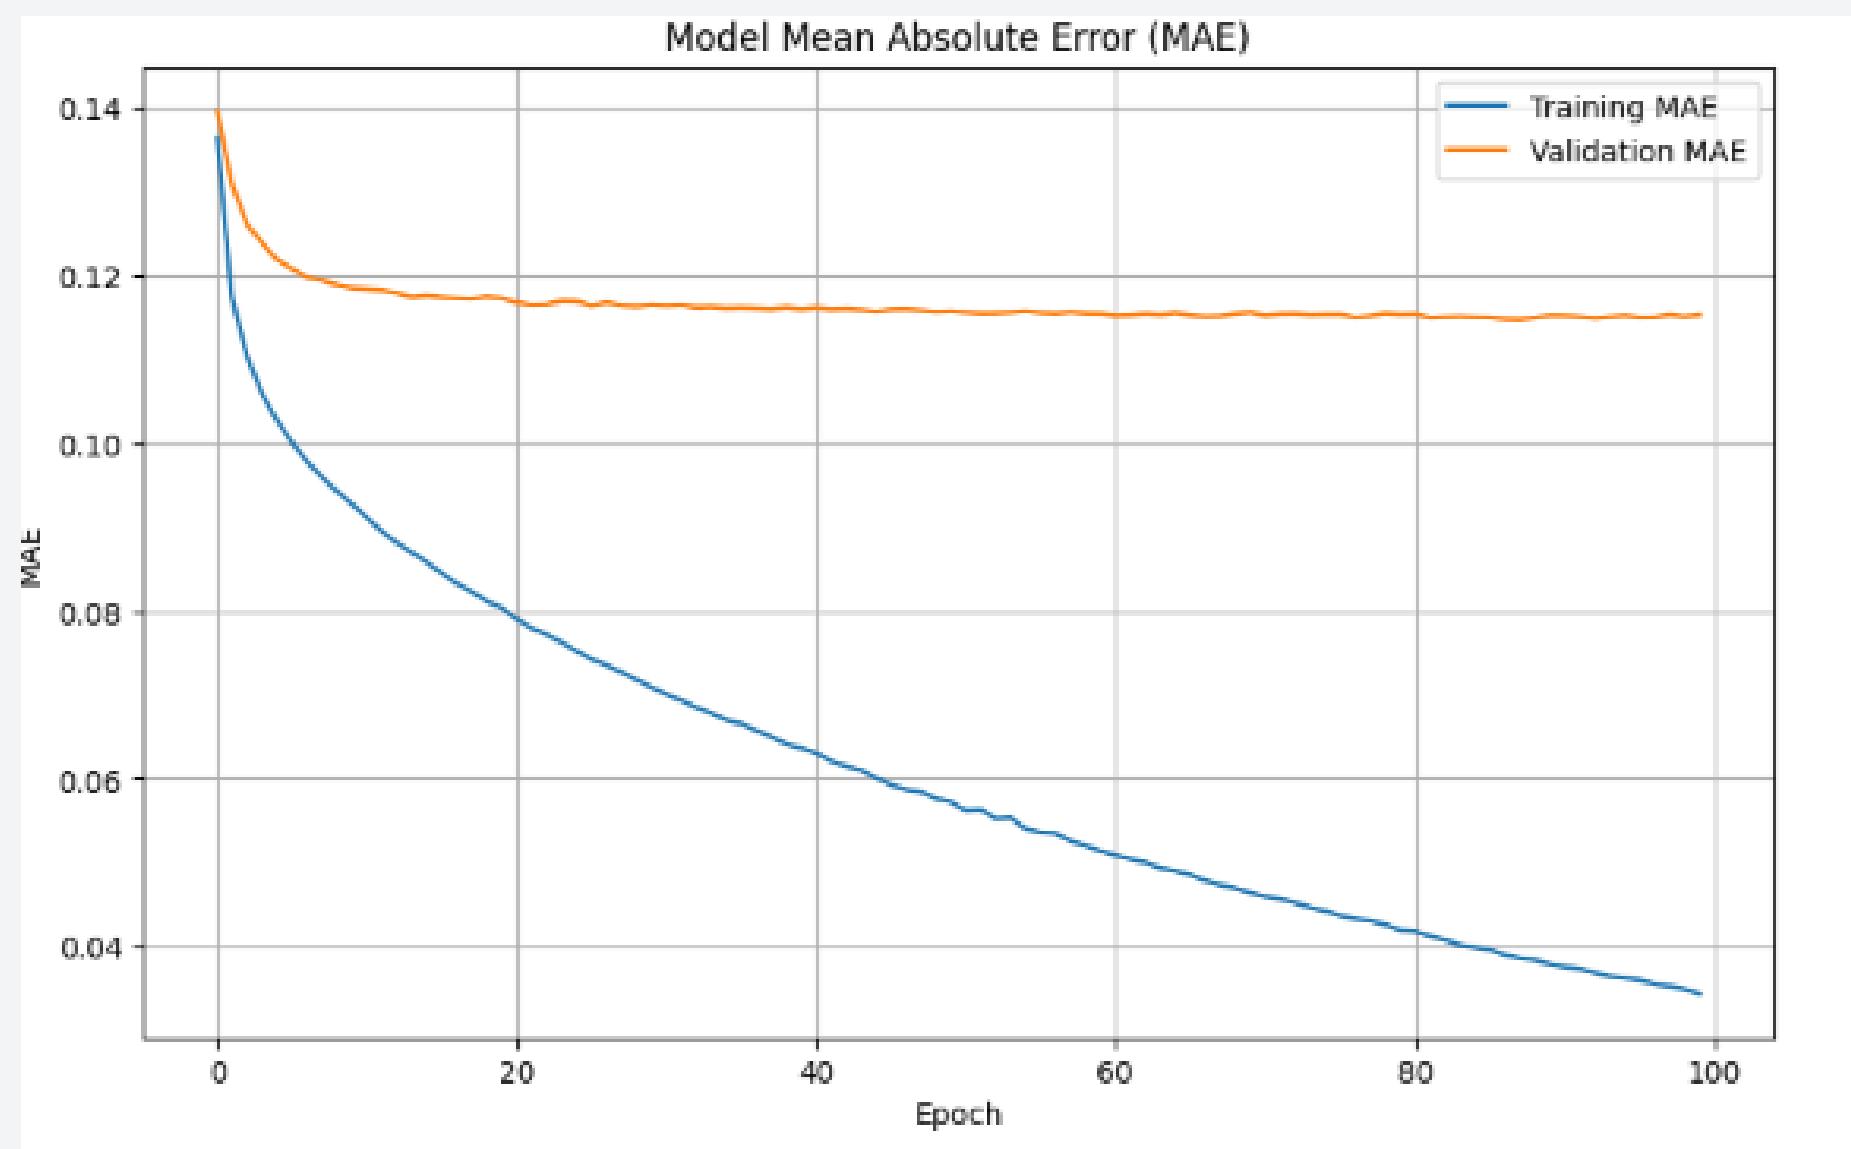
#preprocess_input = tf.keras.applications.mobilenet.preprocess_input
preprocess_input_Mobilenet = tf.keras.applications.mobilenet_v3.preprocess_input

base_model.trainable = False
for layer in base_model.layers[-3:]:
    layer.trainable = True

# Add custom layers
x = tf.keras.layers.GlobalAveragePooling2D()(base_model.output)
x = tf.keras.layers.Dense(1024, activation='relu')(x)
output = tf.keras.layers.Dense(1, activation='sigmoid')(x)

# Create the new model with NASNetMobile as the base
Mobilenet_model = tf.keras.models.Model(inputs=base_model.input, outputs=output)
# Preprocess the data
X_train_processed_Mobilenet = preprocess_input_Mobilenet(np.copy(X_train))
X_valid_processed_Mobilenet = preprocess_input_Mobilenet(np.copy(X_valid))

print("Num GPUs Available: ", len(tf.config.list_physical_devices('GPU')))
```



```
early_stopping_MobileNet = tf.keras.callbacks.EarlyStopping(monitor='val_loss', mode='min', verbose=1, patience=3)

#base_model = tf.keras.applications.ResNet50(input_shape=(224, 224, 3), weights='imagenet', include_top=False)
base_model = tf.keras.applications.MobileNetV3Small(input_shape=(224, 224, 3), weights='imagenet', include_top=False)

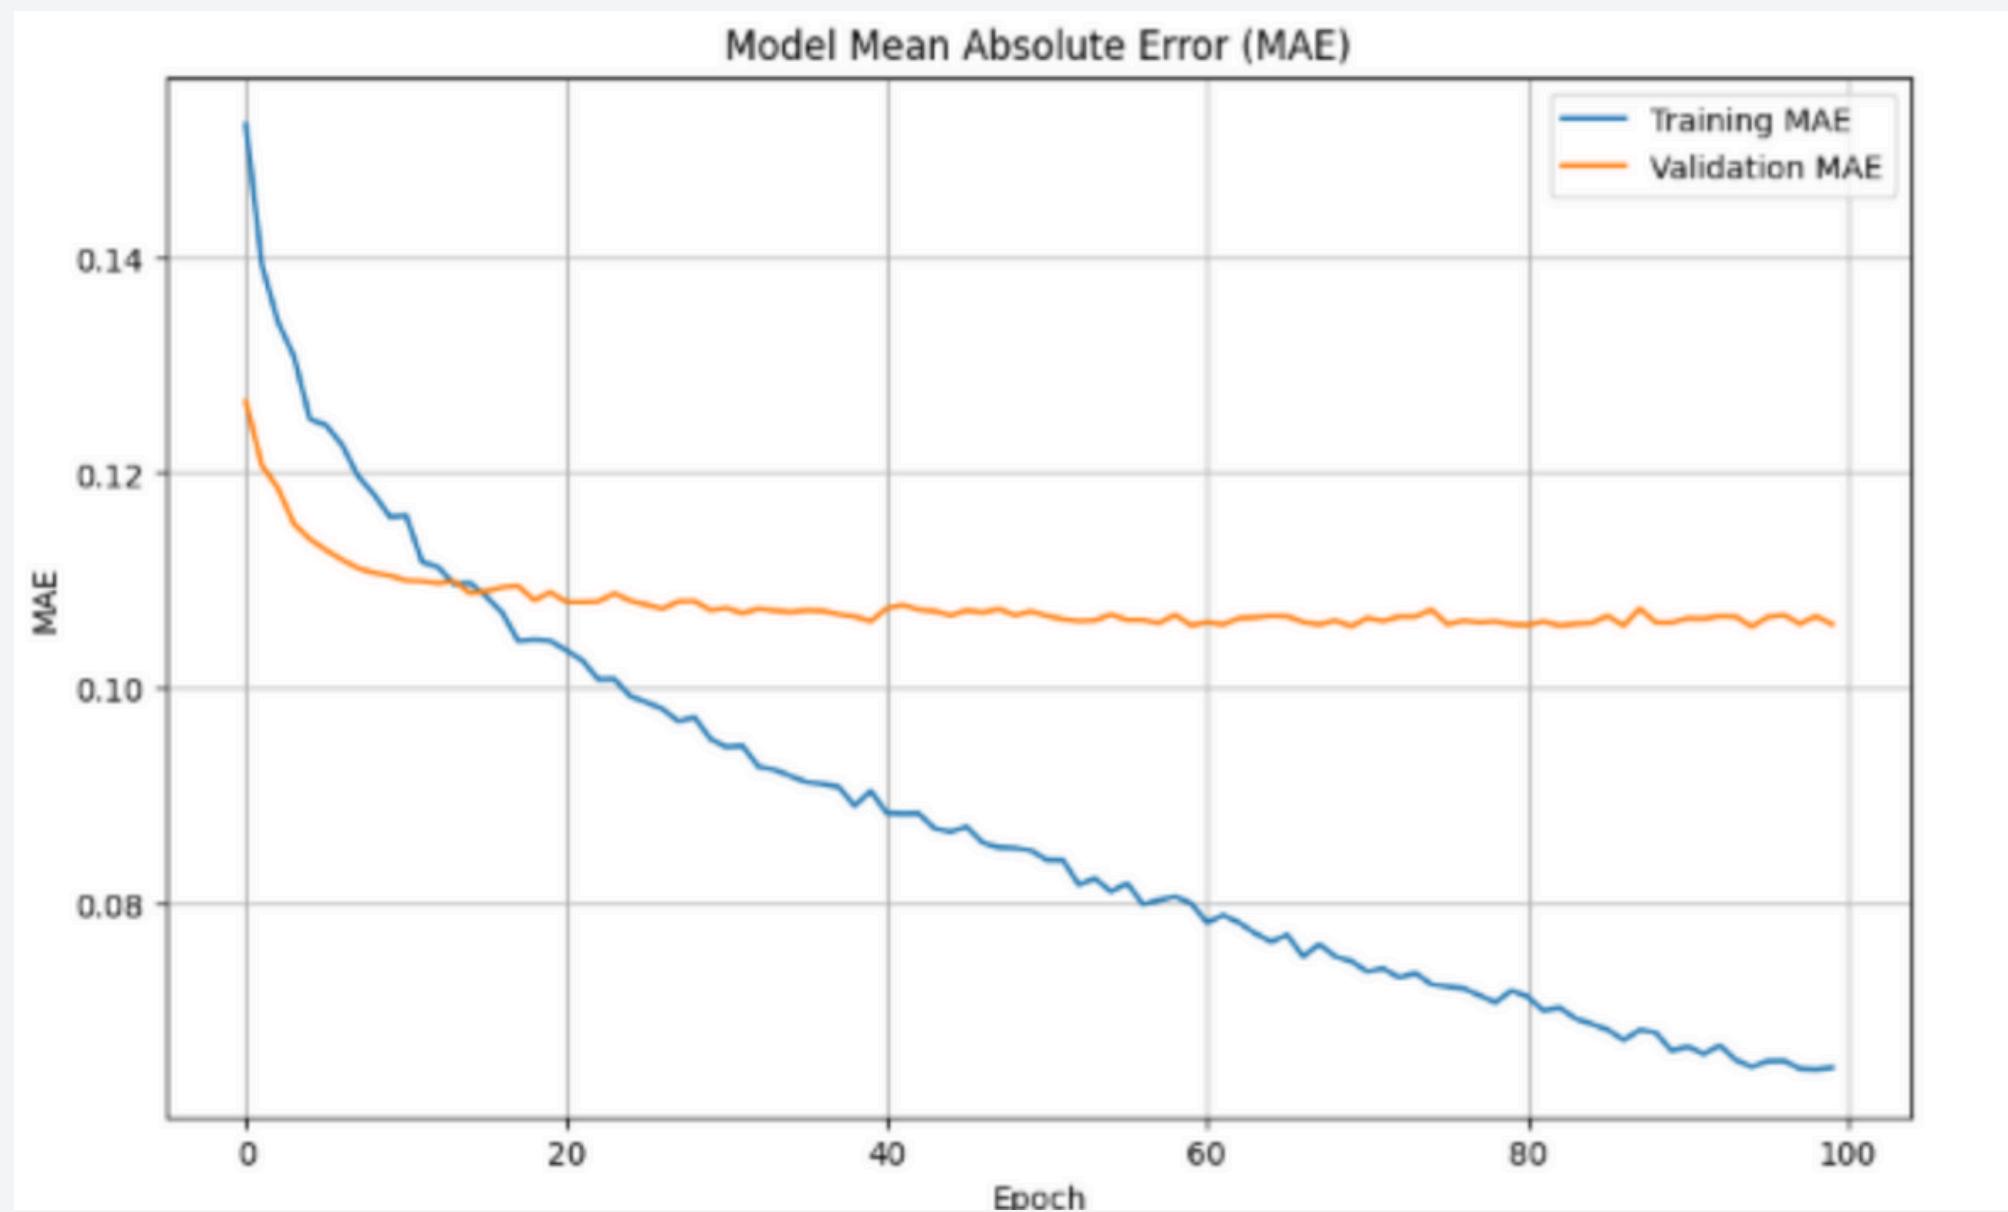
#preprocess_input = tf.keras.applications.mobilenet.preprocess_input
preprocess_input_Mobilenet = tf.keras.applications.mobilenet_v3.preprocess_input

base_model.trainable = False
for layer in base_model.layers[-3:]:
    layer.trainable = True

# Add custom layers
x = tf.keras.layers.GlobalAveragePooling2D()(base_model.output)
x = tf.keras.layers.Dense(1024, activation='relu')(x)
x = tf.keras.layers.Dropout(0.5)(x)
output = tf.keras.layers.Dense(1, activation='sigmoid')(x)

# Create the new model with NASNetMobile as the base
Mobilenet_model = tf.keras.models.Model(inputs=base_model.input, outputs=output)
# Preprocess the data
X_train_processed_Mobilenet = preprocess_input_Mobilenet(np.copy(X_train))
X_valid_processed_Mobilenet = preprocess_input_Mobilenet(np.copy(X_valid))

print("Num GPUs Available: ", len(tf.config.list_physical_devices('GPU')))
```



Train Model
WITH DATA-AUGMENTATION

```
▶ import albumentations as A
from albumentations.core.composition import Compose, OneOf
from albumentations.pytorch import ToTensorV2
import numpy as np
import tensorflow as tf

# Define the Albumentations augmentation pipeline
augmentation_pipeline = A.Compose([
    A.Rotate(limit=10, p=0.5),                      # Rotation range of 10 degrees
    A.HorizontalFlip(p=0.5),                          # Horizontal flip
    A.ShiftScaleRotate(shift_limit=0.1, scale_limit=0.1, rotate_limit=0, p=0.5),
    A.RandomBrightnessContrast(brightness_limit=0.1, contrast_limit=0.1, p=0.5),
    #    A.RandomZoom(scale_limit=0.1, p=0.5),
    A.Normalize(),                                    # Normalization
])
# Custom data generator to apply Albumentations on-the-fly
class AlbumentationsDataGenerator(tf.keras.utils.Sequence):
    def __init__(self, images, labels, batch_size=32, augmentations=None, shuffle=True):
        self.images = images
        self.labels = labels
        self.batch_size = batch_size
        self.augmentations = augmentations
        self.shuffle = shuffle
        self.indices = np.arange(len(self.images))
        self.on_epoch_end()

    def __len__(self):
        return len(self.images) // self.batch_size
```

```
def __getitem__(self, index):
    batch_indices = self.indices[index * self.batch_size:(index + 1) * self.batch_size]
    batch_images = [self.images[i] for i in batch_indices]
    batch_labels = [self.labels[i] for i in batch_indices]

    # Apply augmentations
    if self.augmentations:
        batch_images = [self.augmentations(image=image)['image'] for image in batch_images]
    return np.array(batch_images), np.array(batch_labels)

def on_epoch_end(self):
    if self.shuffle:
        np.random.shuffle(self.indices)

# Initialize the custom data generator with augmentations
train_generator = AlbuumentationsDataGenerator(
    images=X_train,
    labels=Y_train,
    batch_size=32,
    augmentations=augmentation_pipeline
)
Nasnet_model.compile(optimizer=tf.keras.optimizers.Adam(learning_rate=1e-5), loss='mse', metrics=['mae'])

# Training the model
history = Nasnet_model.fit(
    train_generator,
    validation_data=(X_valid_processed_Nasnet, Y_valid),
    epochs=100,
    shuffle=True,
    verbose=1,
    callbacks=[early_stopping_Nasnet]
)
```

✓ 7m

▶ Epoch 12/100
125/125 13s 98ms/step - loss: 0.0269 - mae: 0.1285 - val_loss: 0.0209 - val_mae: 0.1155

→ Epoch 13/100
125/125 20s 95ms/step - loss: 0.0264 - mae: 0.1272 - val_loss: 0.0215 - val_mae: 0.1185

Epoch 14/100
125/125 13s 102ms/step - loss: 0.0247 - mae: 0.1237 - val_loss: 0.0214 - val_mae: 0.1180

Epoch 15/100
125/125 12s 92ms/step - loss: 0.0261 - mae: 0.1259 - val_loss: 0.0210 - val_mae: 0.1166

Epoch 16/100
125/125 12s 94ms/step - loss: 0.0260 - mae: 0.1271 - val_loss: 0.0212 - val_mae: 0.1172

Epoch 17/100
125/125 11s 86ms/step - loss: 0.0242 - mae: 0.1209 - val_loss: 0.0209 - val_mae: 0.1160

Epoch 18/100
125/125 12s 92ms/step - loss: 0.0235 - mae: 0.1199 - val_loss: 0.0209 - val_mae: 0.1162

Epoch 19/100
125/125 11s 88ms/step - loss: 0.0236 - mae: 0.1209 - val_loss: 0.0210 - val_mae: 0.1167

Epoch 20/100
125/125 11s 83ms/step - loss: 0.0236 - mae: 0.1204 - val_loss: 0.0208 - val_mae: 0.1158

Epoch 21/100
125/125 21s 90ms/step - loss: 0.0235 - mae: 0.1202 - val_loss: 0.0207 - val_mae: 0.1156

Epoch 22/100
125/125 11s 88ms/step - loss: 0.0237 - mae: 0.1188 - val_loss: 0.0210 - val_mae: 0.1171

Epoch 23/100
125/125 22s 99ms/step - loss: 0.0220 - mae: 0.1165 - val_loss: 0.0210 - val_mae: 0.1172

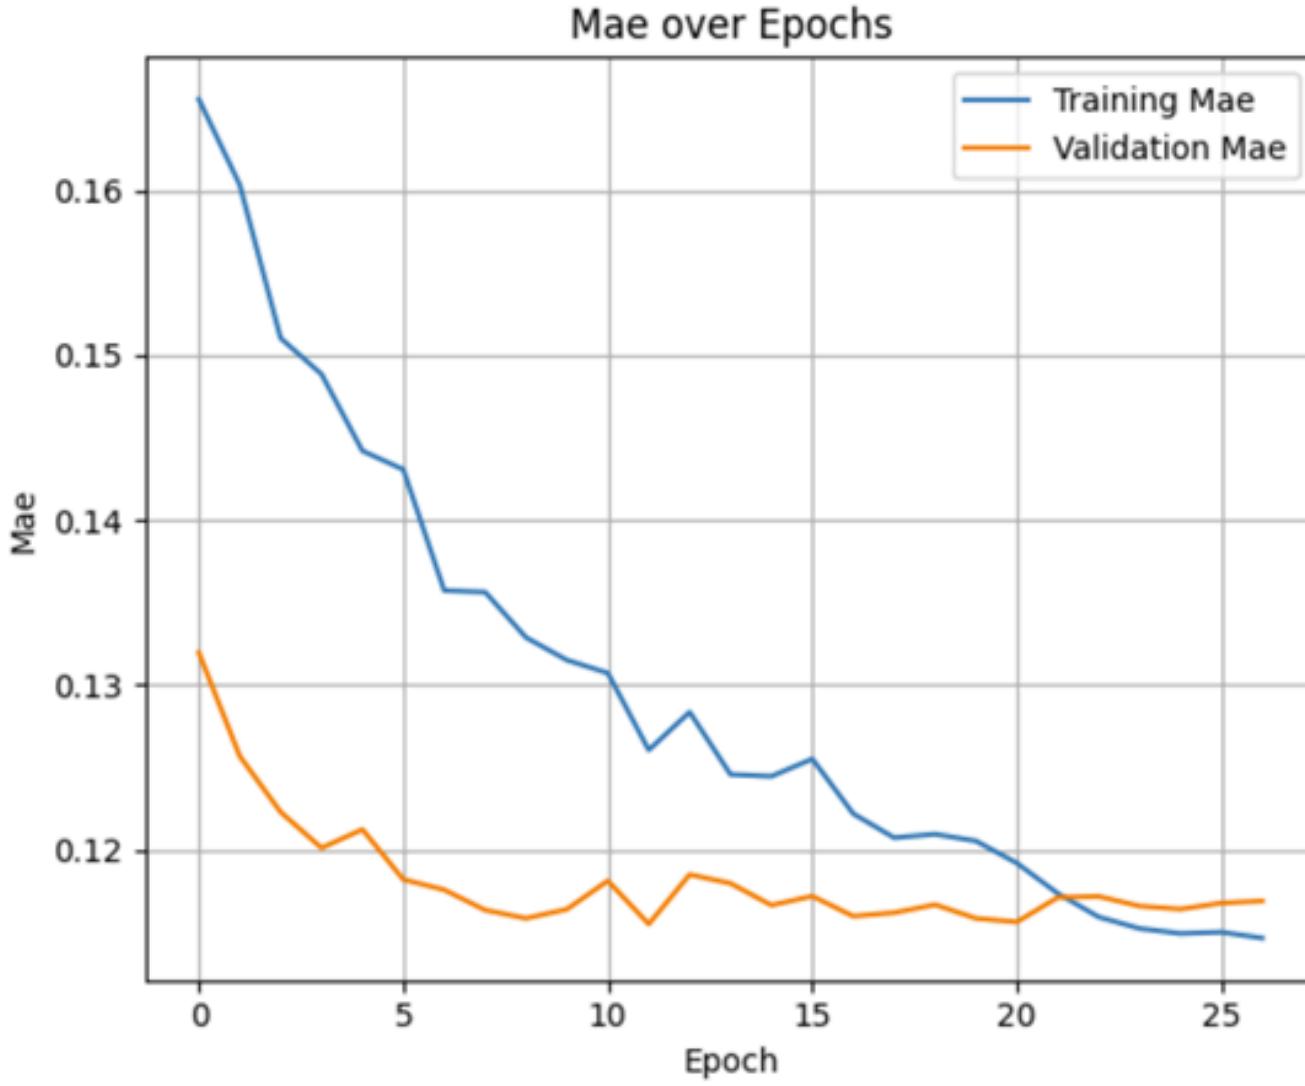
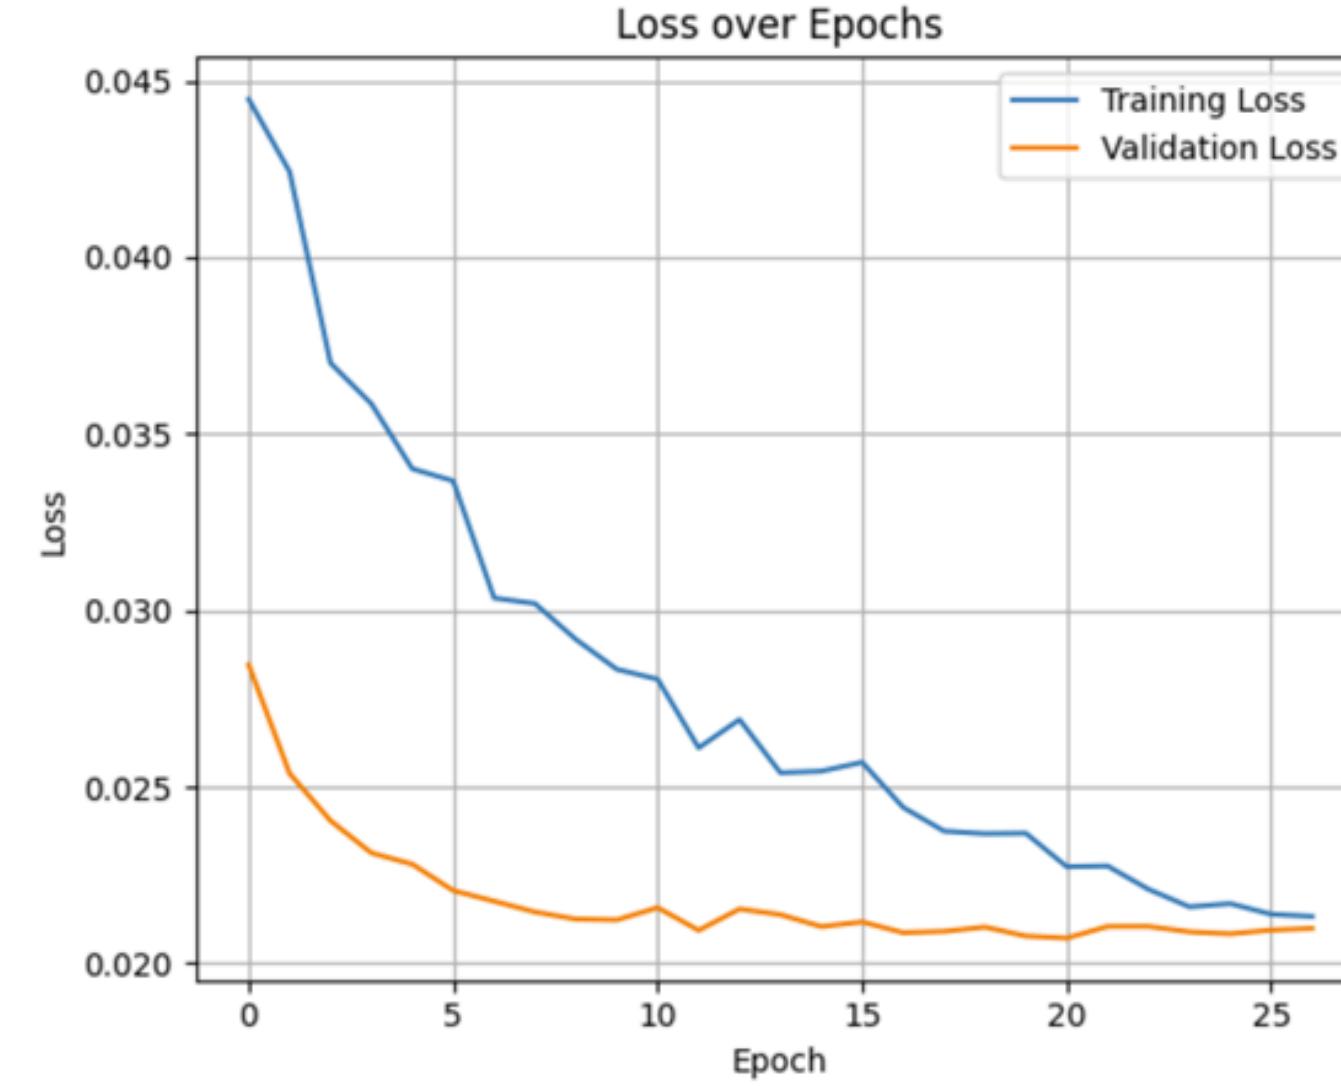
Epoch 24/100
125/125 12s 96ms/step - loss: 0.0218 - mae: 0.1155 - val_loss: 0.0209 - val_mae: 0.1166

Epoch 25/100
125/125 20s 86ms/step - loss: 0.0219 - mae: 0.1159 - val_loss: 0.0208 - val_mae: 0.1164

Epoch 26/100
125/125 21s 90ms/step - loss: 0.0222 - mae: 0.1170 - val_loss: 0.0209 - val_mae: 0.1168

Epoch 27/100
125/125 21s 91ms/step - loss: 0.0208 - mae: 0.1137 - val_loss: 0.0210 - val_mae: 0.1169

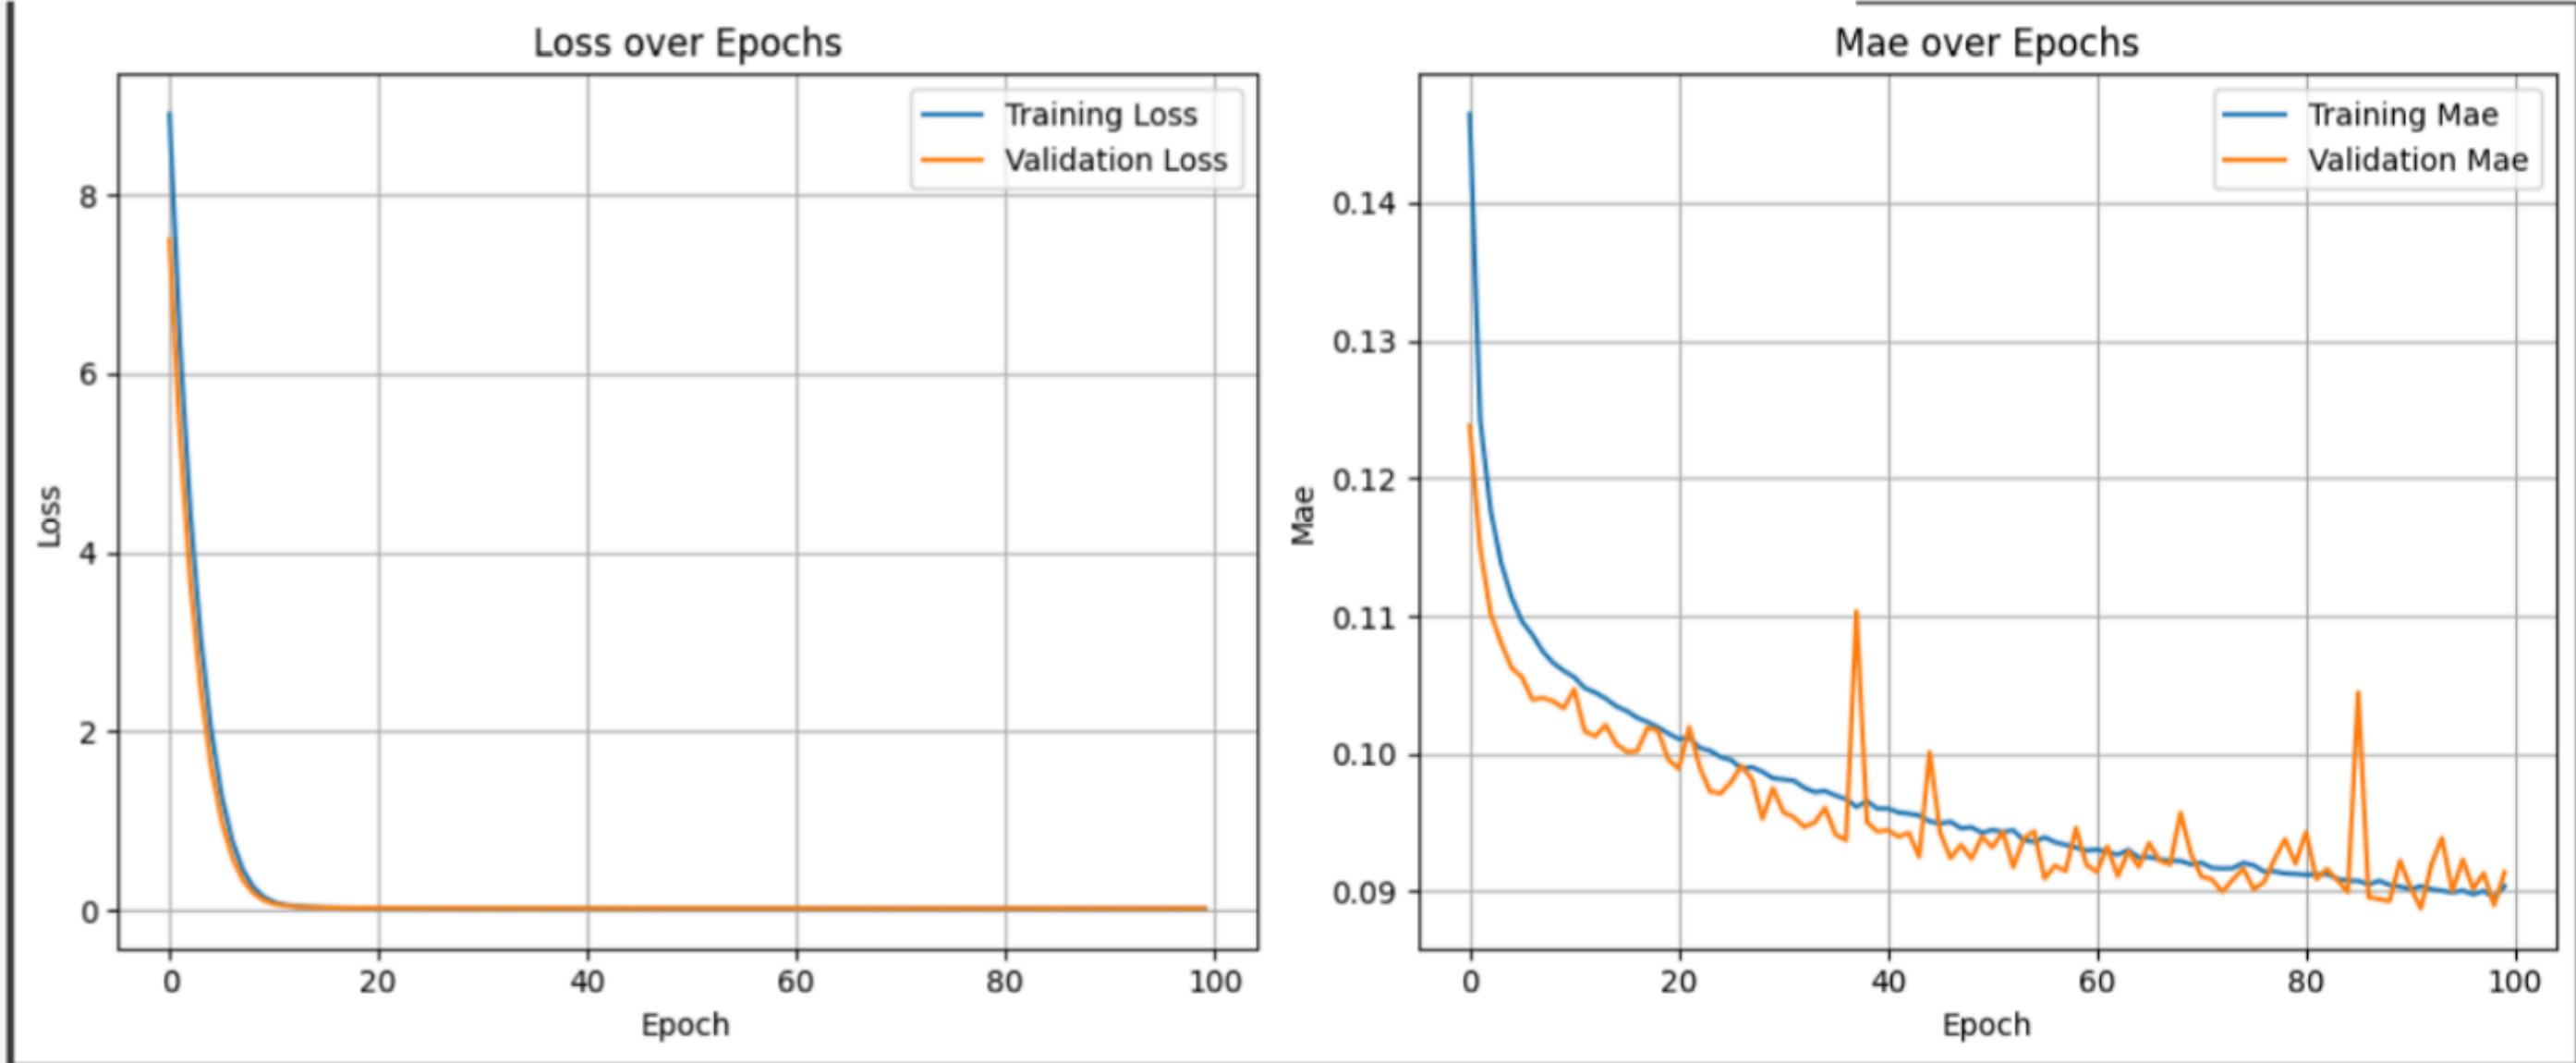
Epoch 27: early stopping



Current learning rate: 1e-05

```
▶ ### Nasnet implementation avec 0.7 de Dropout and with 1 trainable layer
early_stopping_Nasnet = tf.keras.callbacks.EarlyStopping(monitor='val_loss', mode='min', verbose=1, patience=6)
# Load NASNetMobile with pre-trained ImageNet weights, excluding the top layer
base_model = tf.keras.applications.NASNetMobile(input_shape=(224, 224, 3), weights='imagenet', include_top=False)
#preprocess_input = tf.keras.applications.mobilenet.preprocess_input
preprocess_input_Nasnet = tf.keras.applications.nasnet.preprocess_input
base_model.trainable = False
for layer in base_model.layers[-3:]:
    layer.trainable = True
# Adding custom layers
x = tf.keras.layers.GlobalAveragePooling2D()(base_model.output)
x = tf.keras.layers.Dense(1024, activation='relu', kernel_regularizer=tf.keras.regularizers.l2(0.01))(x)
#x = tf.keras.layers.Dropout(0.5)(x) # Dropout with 50% probability
output = tf.keras.layers.Dense(1, activation='sigmoid')(x)
# Create the new model with NASNetMobile as the base
Nasnet_model = tf.keras.models.Model(inputs=base_model.input, outputs=output)
# Preprocess the data
X_train_processed_Nasnet = preprocess_input_Nasnet(np.copy(X_train2))
X_valid_processed_Nasnet = preprocess_input_Nasnet(np.copy(X_valid2))
print("Num GPUs Available: ", len(tf.config.list_physical_devices('GPU')))
print("X_valid_processed_Nasnet shape:", X_valid_processed_Nasnet.shape)
print("Y_valid2 shape:", Y_valid2.shape)
```

4x more examples : 20K



```
    ### Mobilenet basic
    early_stopping_MobileNet = tf.keras.callbacks.EarlyStopping(monitor='val_loss', mode='min', verbose=1, patience=7)

    #base_model = tf.keras.applications.ResNet50(input_shape=(224, 224, 3), weights='imagenet', include_top=False)
    base_model = tf.keras.applications.MobileNetV3Small(input_shape=(224, 224, 3), weights='imagenet', include_top=False)

    #preprocess_input = tf.keras.applications.mobilenet.preprocess_input
    preprocess_input_Mobilenet = tf.keras.applications.mobilenet_v3.preprocess_input

    base_model.trainable = False

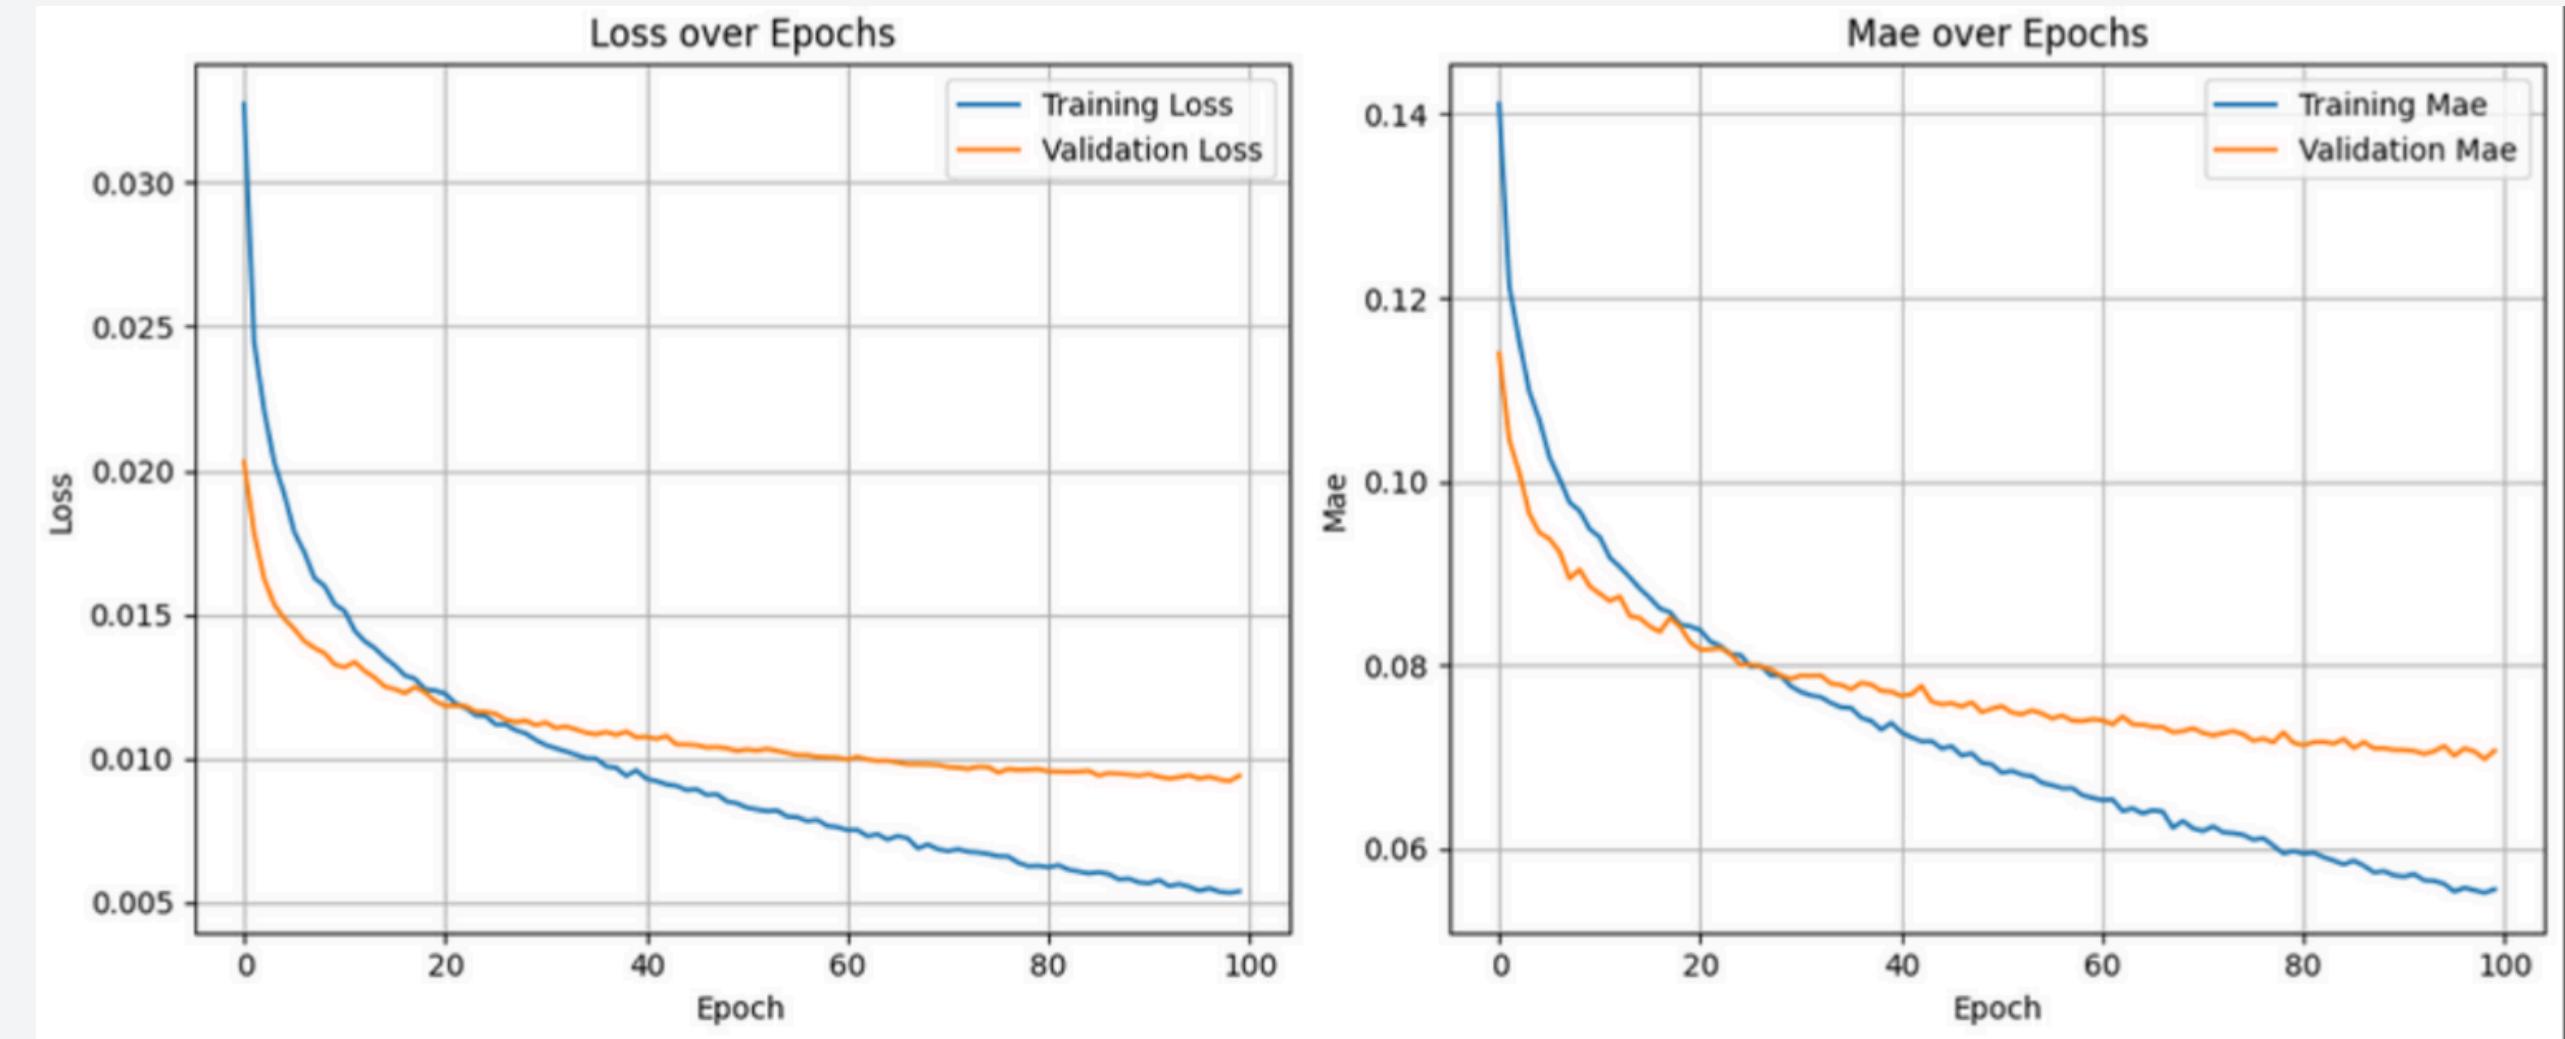
    base_model.trainable = False
    for layer in base_model.layers[-3:]:
        layer.trainable = True

    # Add custom layers
    x = tf.keras.layers.GlobalAveragePooling2D()(base_model.output)
    x = tf.keras.layers.Dense(1024, activation='relu')(x)
    # = tf.keras.layers.Dropout(0.5)(x)
    output = tf.keras.layers.Dense(1, activation='sigmoid')(x)

    # Create the new model with MobileNet as the base
    Mobilenet_model = tf.keras.models.Model(inputs=base_model.input, outputs=output)
    # Preprocess the data
    X_train_processed_Mobilenet = preprocess_input_Mobilenet(np.copy(X_train2))
    X_valid_processed_Mobilenet = preprocess_input_Mobilenet(np.copy(X_valid2))

    print("Num GPUs Available: ", len(tf.config.list_physical_devices('GPU')))

    print(f'Nb layer : {len(Mobilenet_model.layers)}')
```



Train Model

MODEL ENSEMBLING

```
✓ [81] from sklearn.metrics import mean_absolute_error
    Nasnet_prediction = Nasnet_model.predict(X_valid2)
    Mobilenet_prediction = Mobilenet_model.predict(X_valid2)
    predictions = (Nasnet_prediction*0.4) + (0.6*Mobilenet_prediction)

→ 32/32 ━━━━━━━━━━ 1s 20ms/step
32/32 ━━━━━━━━━━ 0s 5ms/step

✓ 0s
    mae = mean_absolute_error(Y_valid2, predictions)
    print(f"Mean Absolute Error (MAE): {mae}")

→ Mean Absolute Error (MAE): 0.8256474798896908
```

Train Model

CLASSIFICATION

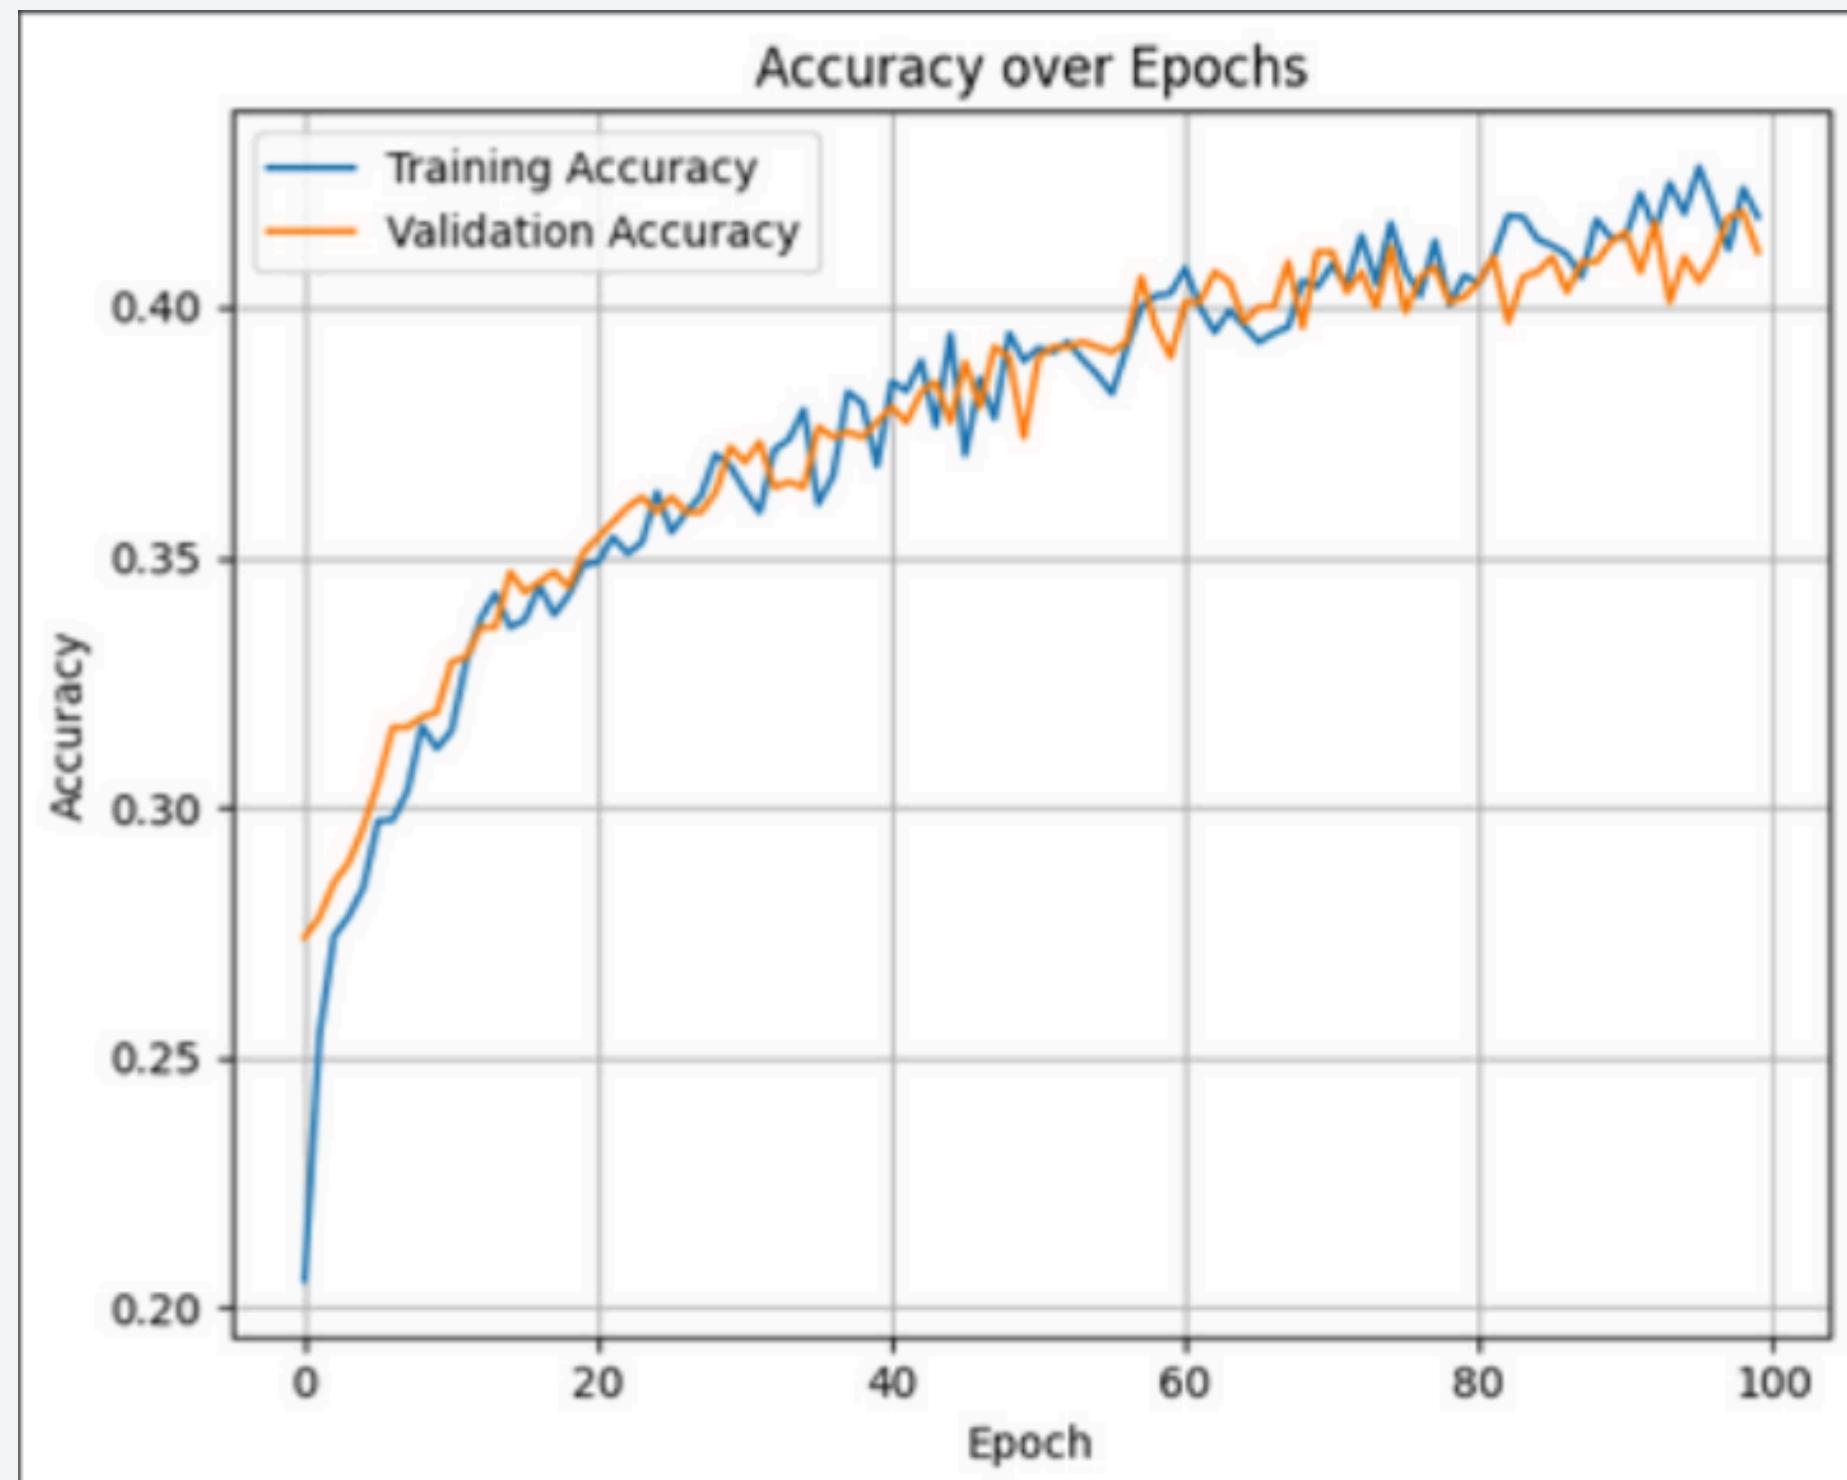
Age Buckets

```
[13] # defining category

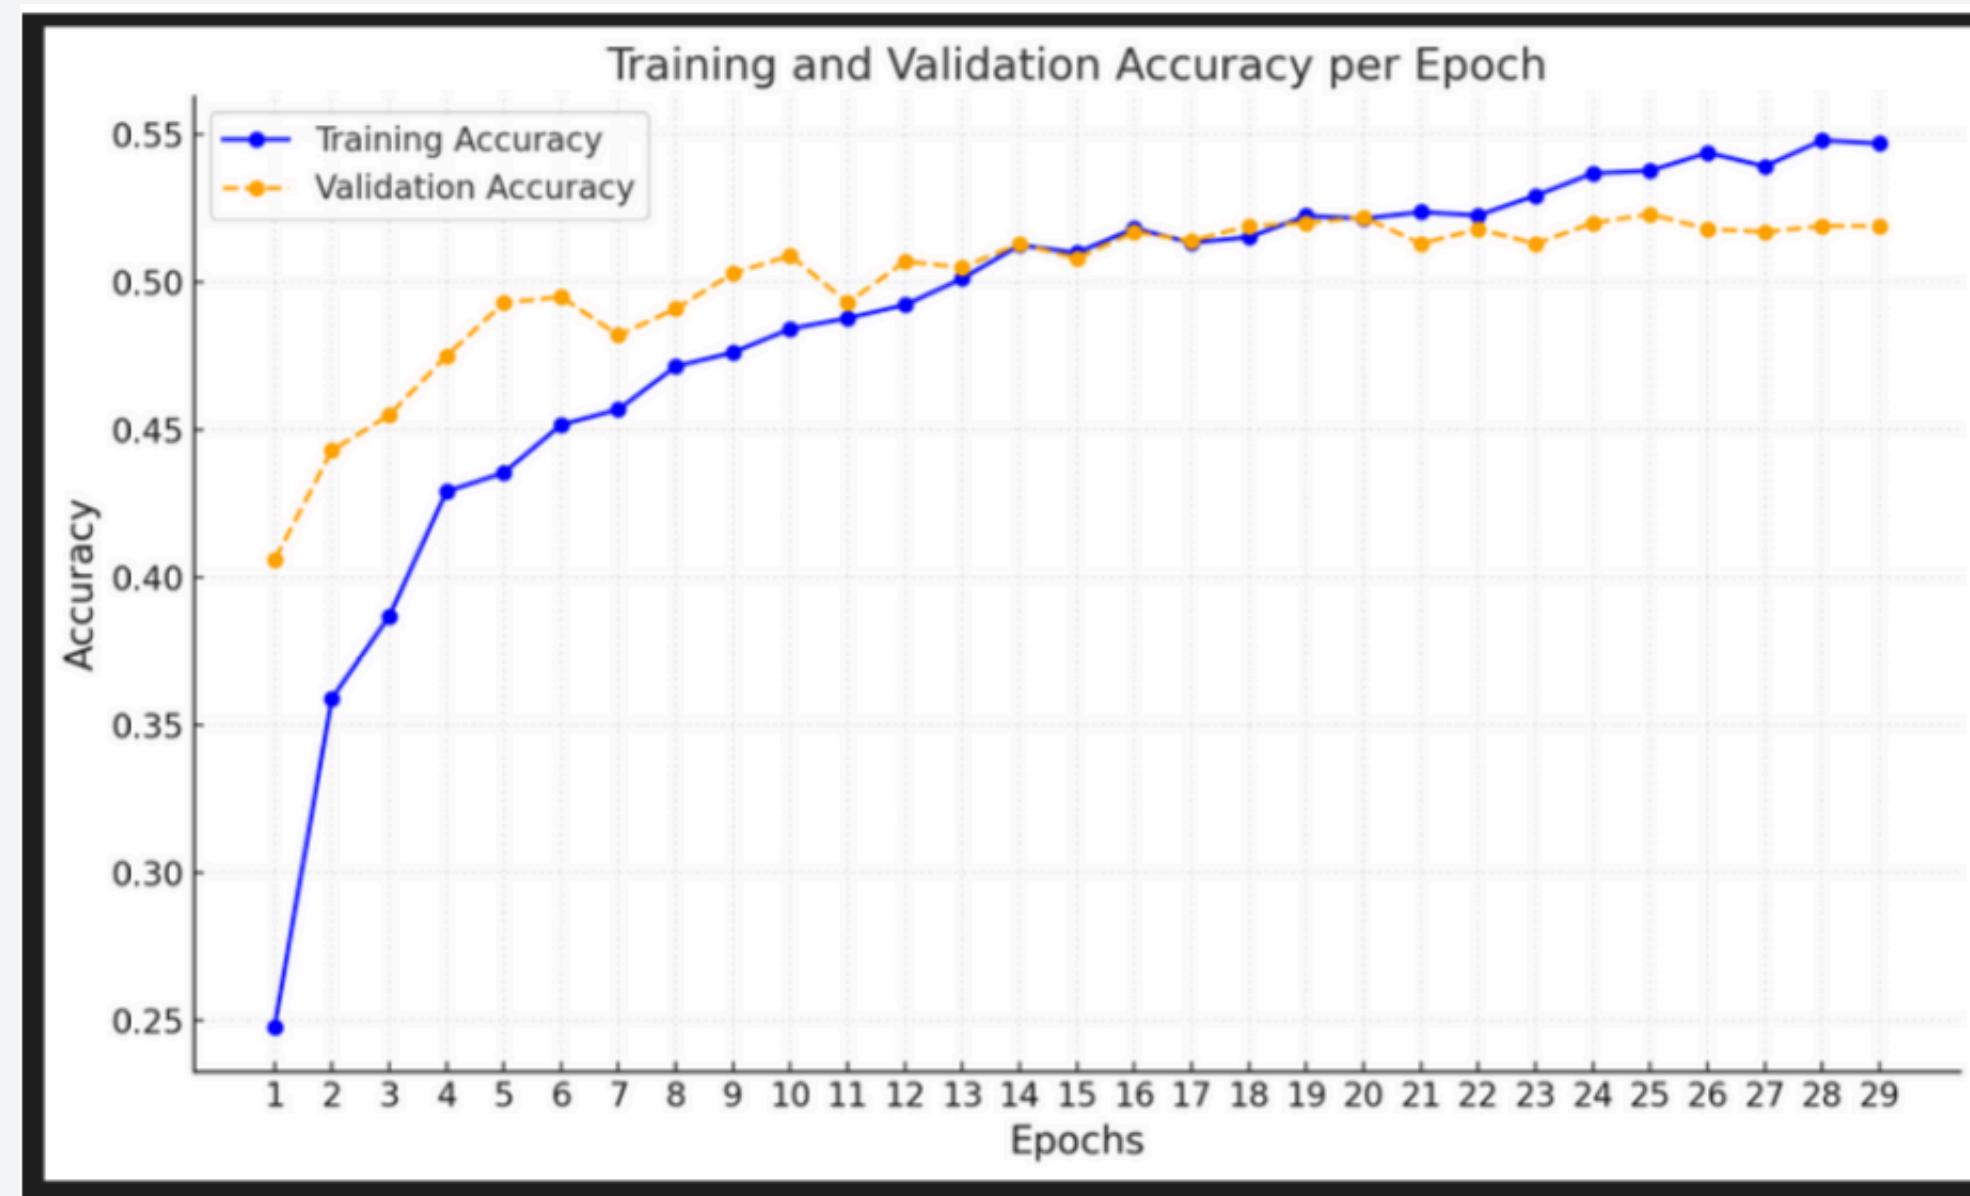
def age_to_class(age):
    age = int(age)
    if 0 <= age <= 2:
        return 0
    elif 3 <= age <= 9:
        return 1
    elif 10 <= age <= 20:
        return 2
    elif 21 <= age <= 27:
        return 3
    elif 28 <= age <= 45:
        return 4
    elif 46 <= age <= 65:
        return 5
    else:
        return 6
```

1 sigmoid -> 7 softmax

NasNet



MobileNet V3 Small



Comparison

with

Production Level Models

Model	Dataset	MAE (Years)	Accuracy (%)
VGGNet	IMDB-WIKI	~4.2–4.8	~75–80%
	FG-NET	~3.1–3.5	~80–85%
ResNet	IMDB-WIKI	~3.2–3.8	~82–88%
	FG-NET	~2.8–3.0	~85–90%
	APPA-REAL	~3.5–4.0	~78–83%
EfficientNet	IMDB-WIKI	~2.8–3.5	~85–90%
	FG-NET	~2.5–2.8	~88–92%
	APPA-REAL	~3.2–3.8	~82–87%
Vision Transformers (ViT)	IMDB-WIKI	~2.6–3.2	~87–92%
	FG-NET	~2.4–2.7	~90–94%
	APPA-REAL	~3.0–3.5	~85–90%

Nasnet (TL)

val_mae: ~9

accuracy: ~45%

MobileNet (TL)

val_mae: ~7

accuracy: ~55%

Thank you!